

# Administrivia

## Updated Midterm/Final grading policy

- if your final has a higher score, we will drop your midterm I score and weigh the final represent both exams

# Administrivia

## HW4

- Query optimization
- EXPLAIN query

## Project 2

- Graph analysis using SQL
- Extra Credit: recursive WITH clauses

Both Due 11/21

# Real-world Example of FDs

## Airport Dataset

```
"id","ident","type","name","latitude_deg","longitude_deg","elevation_ft","continent","iso_country","iso_region","municipality","sche  
6523,"00A","heliport","Total Rf Heliport",40.07080078125,-74.93360137939453,11,"NA","US","US-PA","Bensalem","no","00A","","00A",,,  
323361,"00AA","small_airport","Aero B Ranch Airport",38.704022,-101.473911,3435,"NA","US","US-KS","Leoti","no","00AA","","00AA",,,  
6524,"00AK","small_airport","Lowell Field",59.947733,-151.692524,450,"NA","US","US-AK","Anchor Point","no","00AK","","00AK",,,  
6525,"00AL","small_airport","Epps Airpark",34.86479949951172,-86.77030181884766,820,"NA","US","US-AL","Harvest","no","00AL","","00AL",,  
6526,"00AR","closed","Newport Hospital & Clinic Heliport",35.6087,-91.254898,237,"NA","US","US-AR","Newport","no",,,,,"00AR"  
322127,"00AS","small_airport","Fulton Airport",34.9428028,-97.8180194,1100,"NA","US","US-OK","Alex","no","00AS","","00AS",,,  
6527,"00AZ","small_airport","Cordes Airport",34.305599212646484,-112.16500091552734,3810,"NA","US","US-AZ","Cordes","no","00AZ","","00AZ",,,  
6528,"00CA","small_airport","Goldstone (GTS) Airport",35.35474,-116.885329,3038,"NA","US","US-CA","Barstow","no","00CA","","00CA",,,  
324424,"00CL","small_airport","Williams Ag Airport",39.427188,-121.763427,87,"NA","US","US-CA","Biggs","no","00CL","","00CL",,,  
322658,"00CN","heliport","Kitchen Creek Helibase Heliport",32.7273736,-116.4597417,3350,"NA","US","US-CA","Pine Valley","no","00CN",  
6529,"00CO","closed","Cass Field",40.622202,-104.344002,4830,"NA","US","US-CO","Briggsdale","no",,,,,"00CO"  
6531,"00FA","small_airport","Grass Patch Airport",28.64550018310547,-82.21900177001953,53,"NA","US","US-FL","Bushnell","no","00FA",,  
6532,"00FD","closed","Ringhaver Heliport",28.8466,-82.345398,25,"NA","US","US-FL","Riverview","no",,,,,"00FD"  
6533,"00FL","small_airport","River Oak Airport",27.230899810791016,-80.96920013427734,35,"NA","US","US-FL","Okeechobee","no","00FL",  
534,"00GA","small_airport","Lt World Airport",33.76750183105469,-84.06829833984375,700,"NA","US","US-GA","Lithonia","no","00GA",,,  
6535,"00GE","heliport","Caffrey Heliport",33.887982,-84.736983,957,"NA","US","US-GA","Hiram","no","00GE","","00GE",,,  
6536,"00HI","heliport","Kaupulehu Heliport",19.832881,-155.978347,43,"OC","US","US-HI","Kailua-Kona","no","00HI","","00HI",,,
```

# TPC-H Warehousing Benchmark

LineID, OrderID, PartID, SupplierID, Qty, Discount, Tax,  
CustomerID, OrderTotalPrice, OrderDate,  
AvailQty, SupplyCost,  
PartName, Brand, Type  
SupplierName, Address

Order → CustomerID, OrderTotalPrice, OrderDate

PartID,SupplierID → AvailQty, SupplyCost

PartID → PartName, Brand, Type

SupplierID → SupplierName, Address

LineID → OrderID, PartID, SupplierID, Qty, Discount, Tax

# Disk, Storage, and Indexing

Eugene Wu

# DBMS Overview

Each layer provides a simple abstraction to layers above it, and makes assumptions about layers below it.

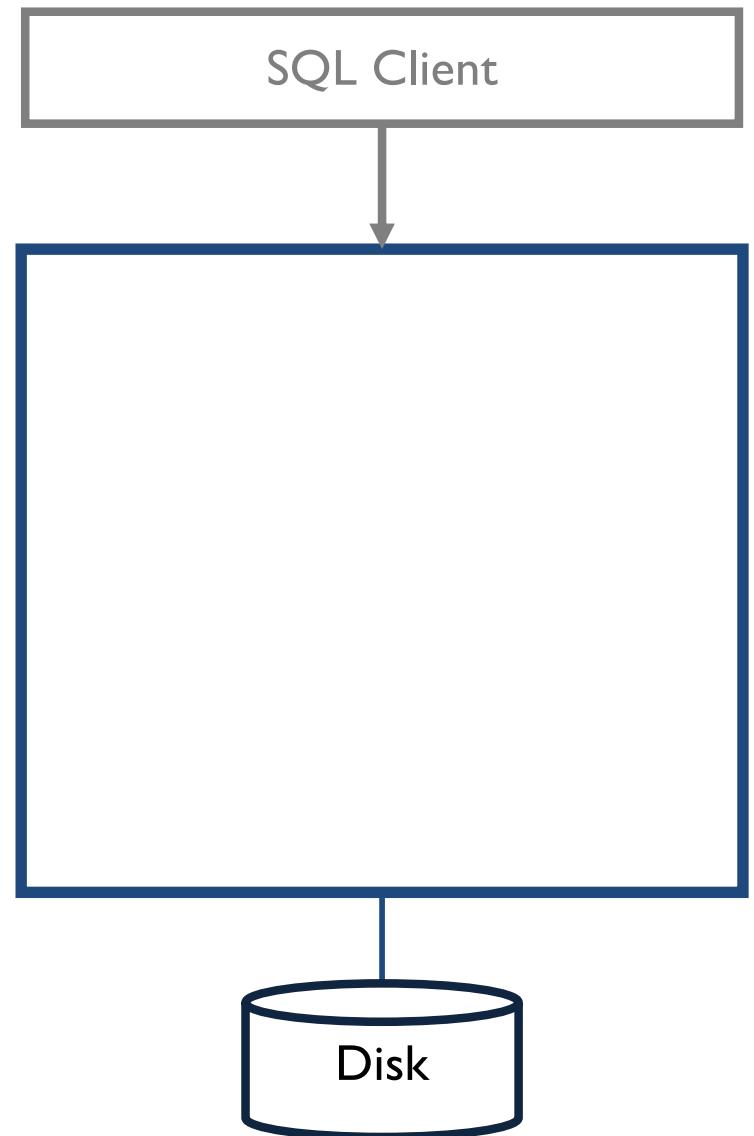
Requires careful design and assumptions for performance



# DBMS Overview

Apps interact w/ SQL Client

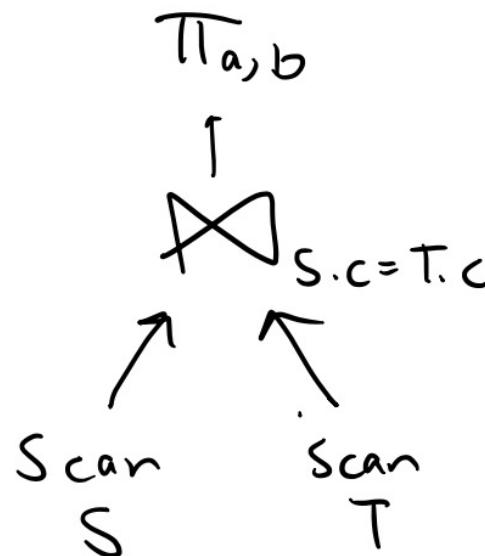
```
db.execute(''  
    SELECT a, b  
        FROM S, T  
    WHERE S.c = T.c  
    ''')
```



# DBMS Overview

Parse, check, & verify SQL query

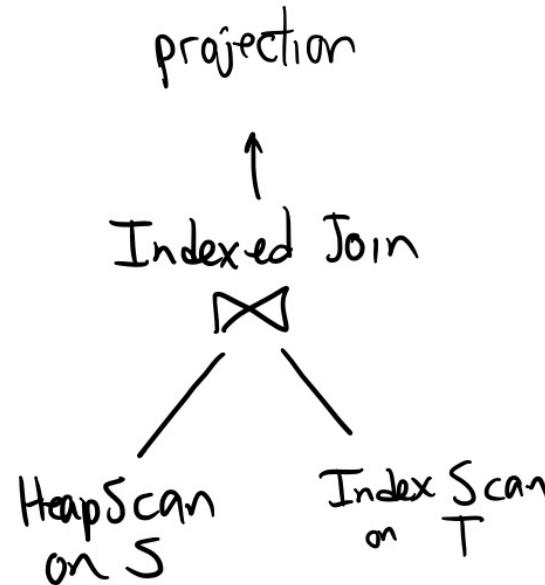
Turn into efficient query plan composed of logical operators



# DBMS Overview

Query is modeled as a relational data flow

Each operator is a specific implementation

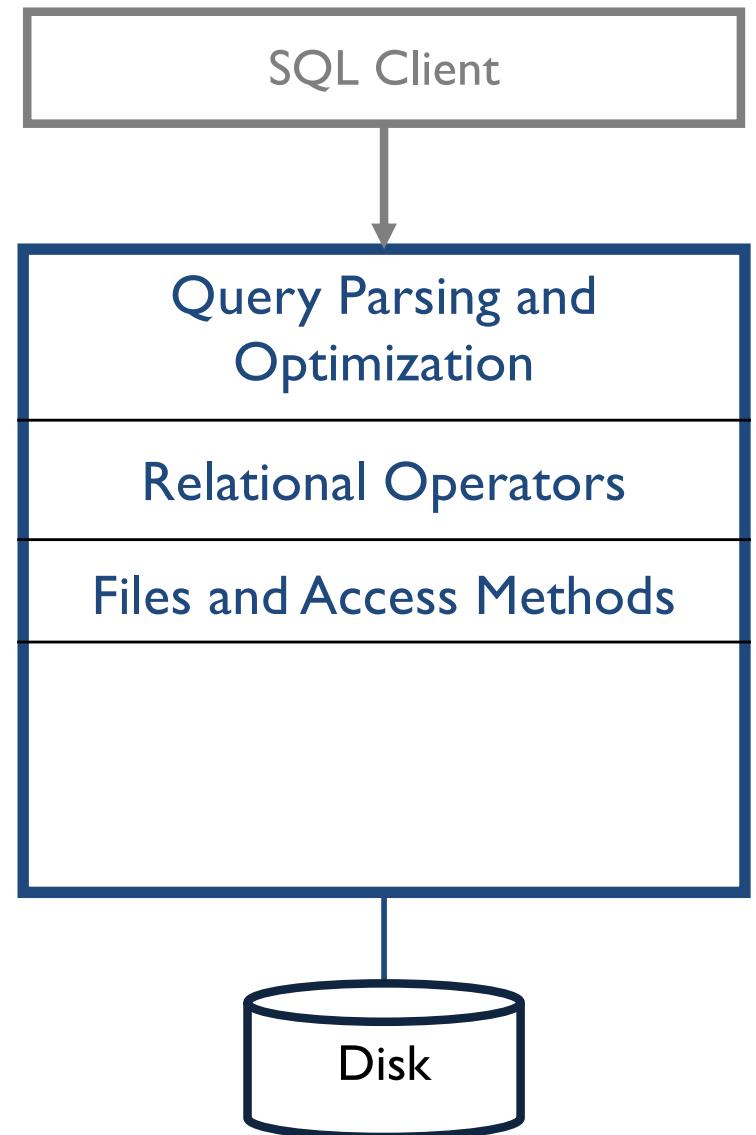


# DBMS Overview

Organizes tables, indexes, records as groups of pages in a “logical file”

API:

- Operators ask for records
- Logical files help read and write bytes on pages



# DBMS Overview

Not all pages can fit into RAM.

Buffer manager provides illusion that all pages are accessible.

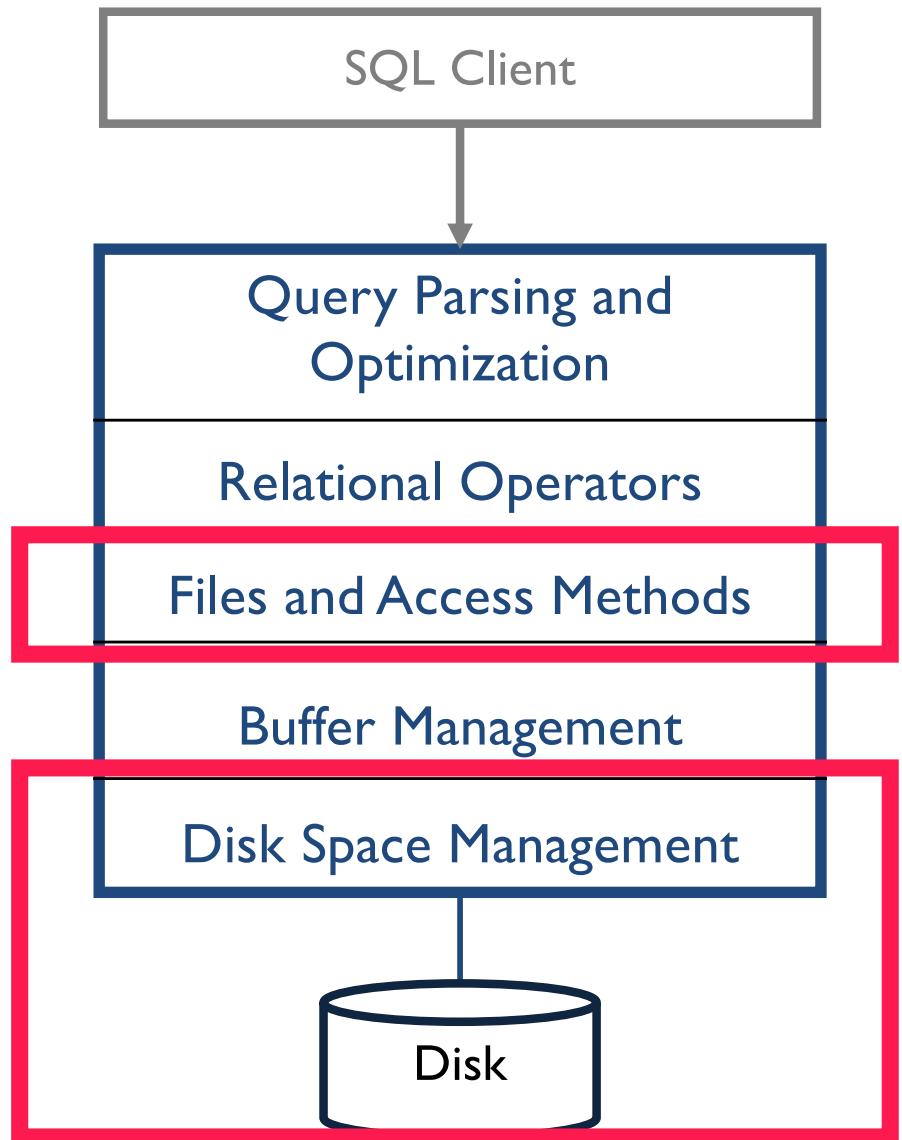
Files simply ask for pages.



# DBMS Overview

Physically read and write bytes on one or more storage devices  
(hard drives, SSDs, etc)

Storage performance properties dictate the design of layers above.



# DBMS Overview

Layers help manage engineering complexity.

Requires assumptions about performance of lower layers. (cost models)



# Storage Devices

Hard Drives vs SSDs vs RAM

Locality (random vs sequential)

Reads vs Writes

Performance and Monetary Costs

# Why not store all in RAM? \$\$\$

Too much \$\$\$

High-end DBs: ~Petabyte (1000TB).

SQL Hyperscale: 100TB+ (2018)

Disks are ~60% cost of a production system

Main memory not persistent

Obviously important if DB stops/crashes

*main-memory* DBMSes discussed in advanced DB course

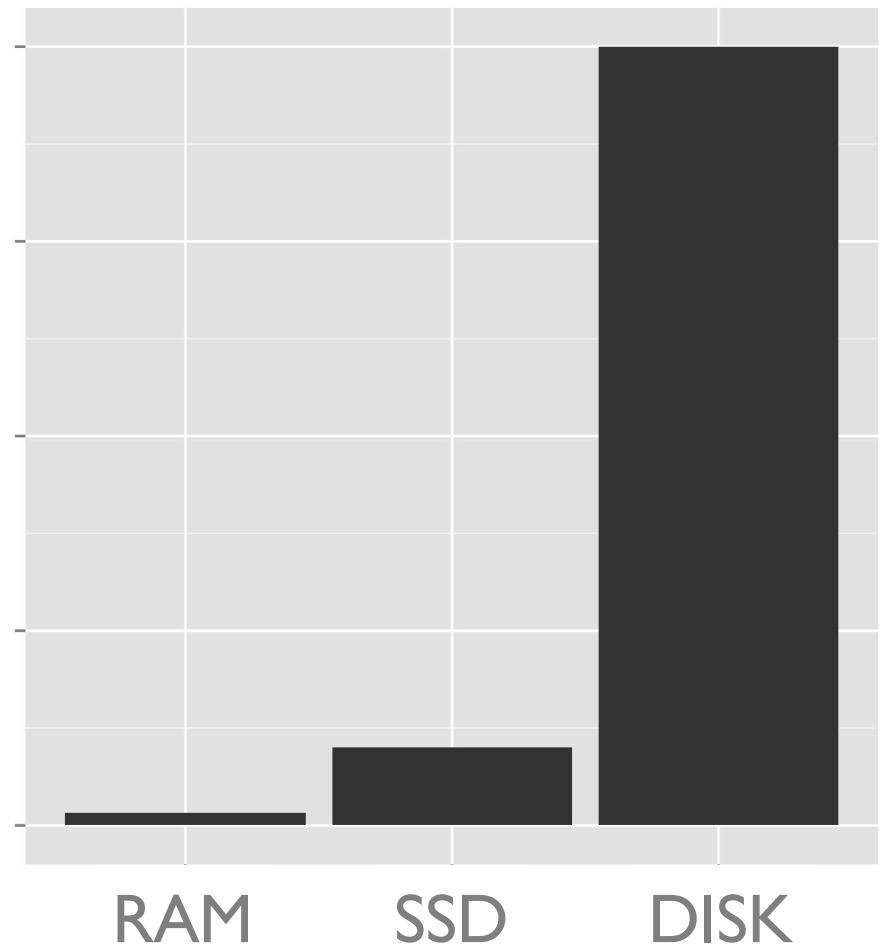
# \$ Matters

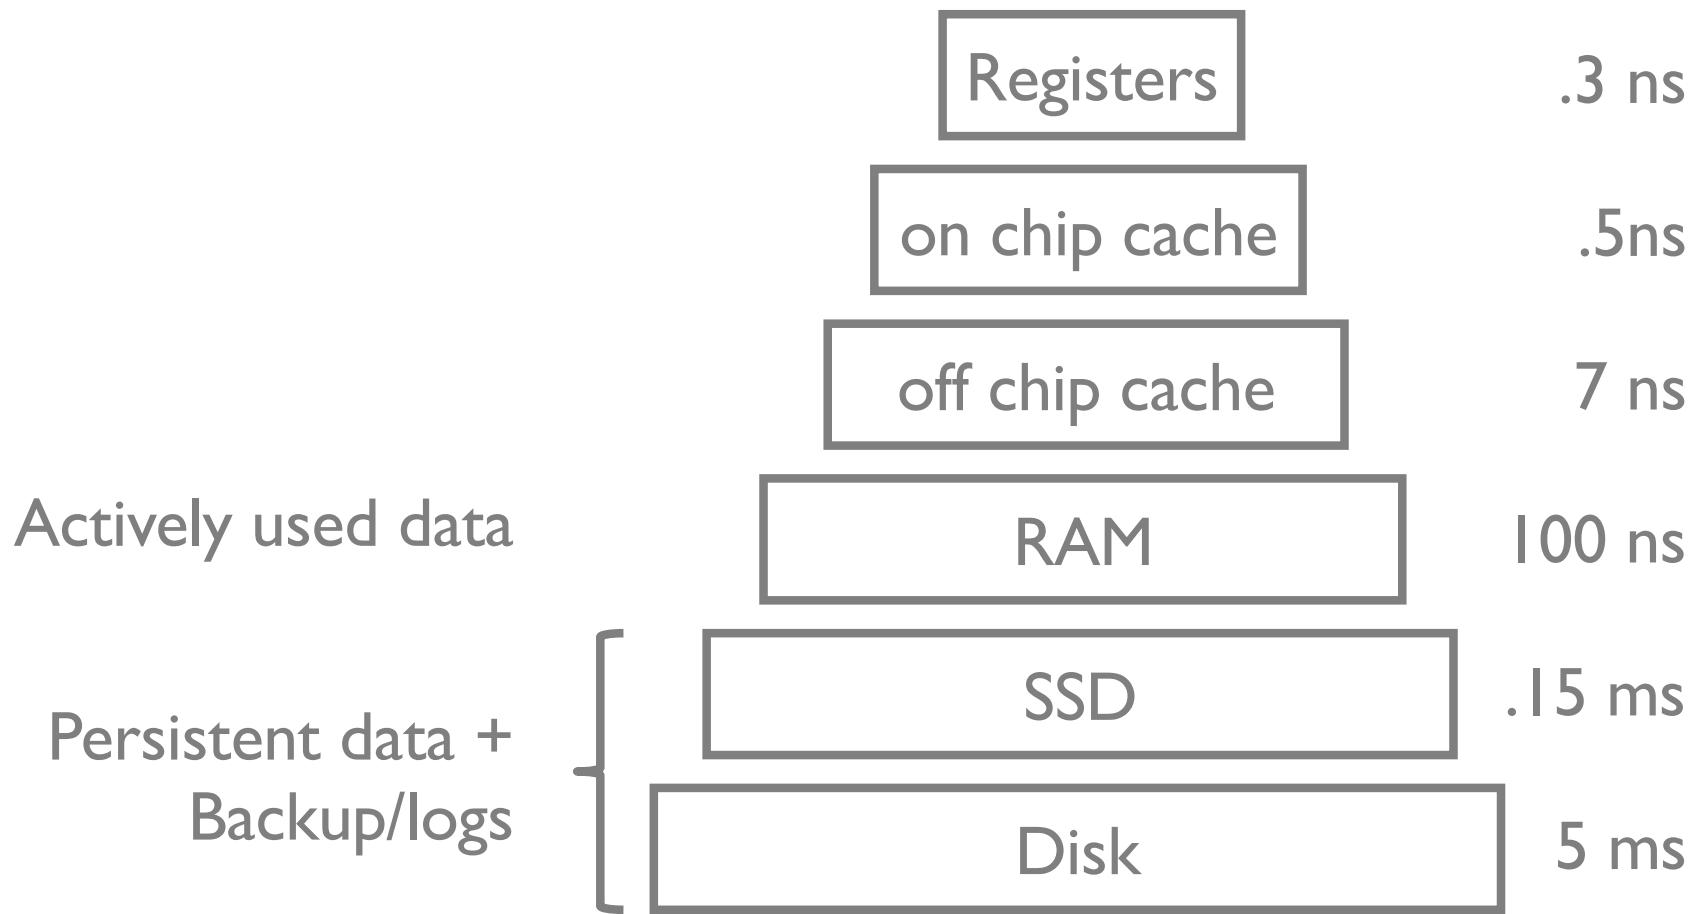
Newegg enterprise \$1000

RAM: 0.08TB

SSD: ~2TB (25x)

Disk: ~40TB (500x)

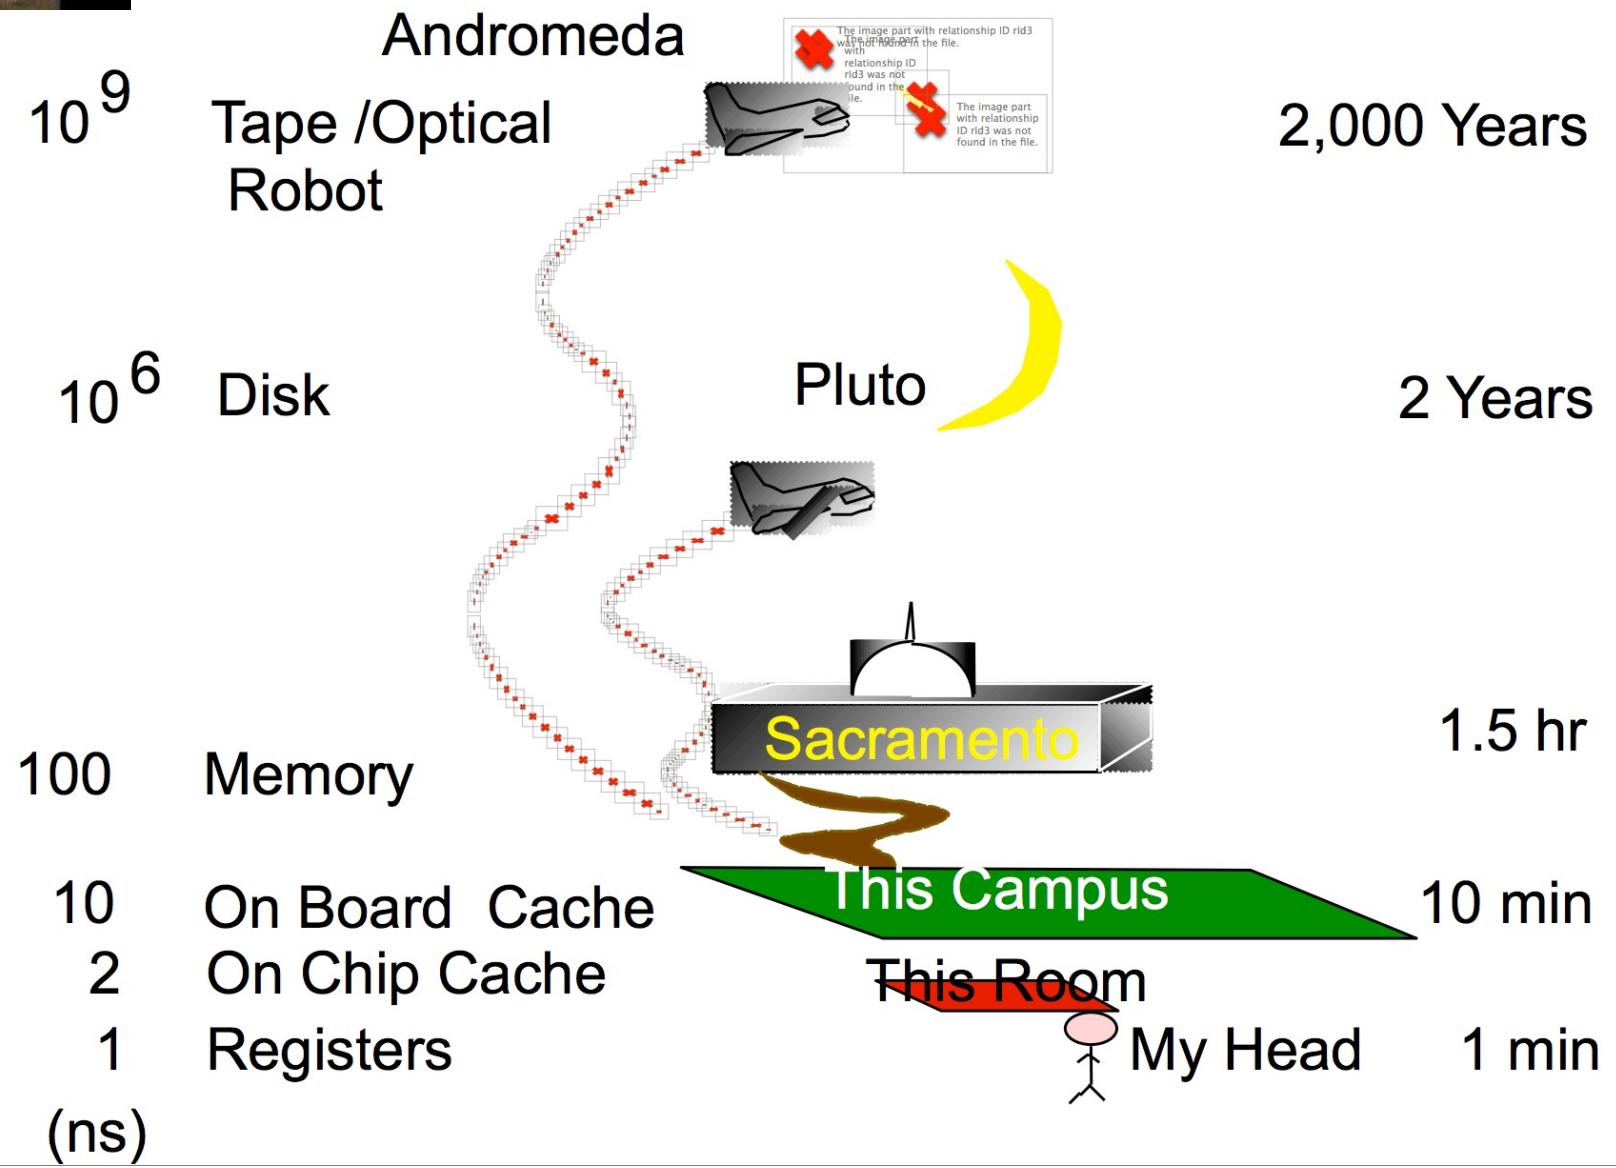
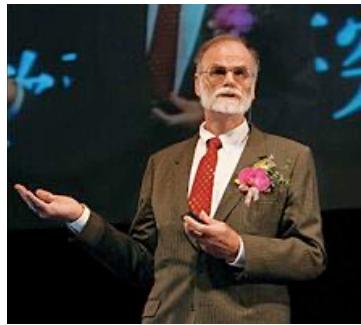




### Interesting numbers

compress 1k bytes: 3000 ns

roundtrip in data center: .5 ms

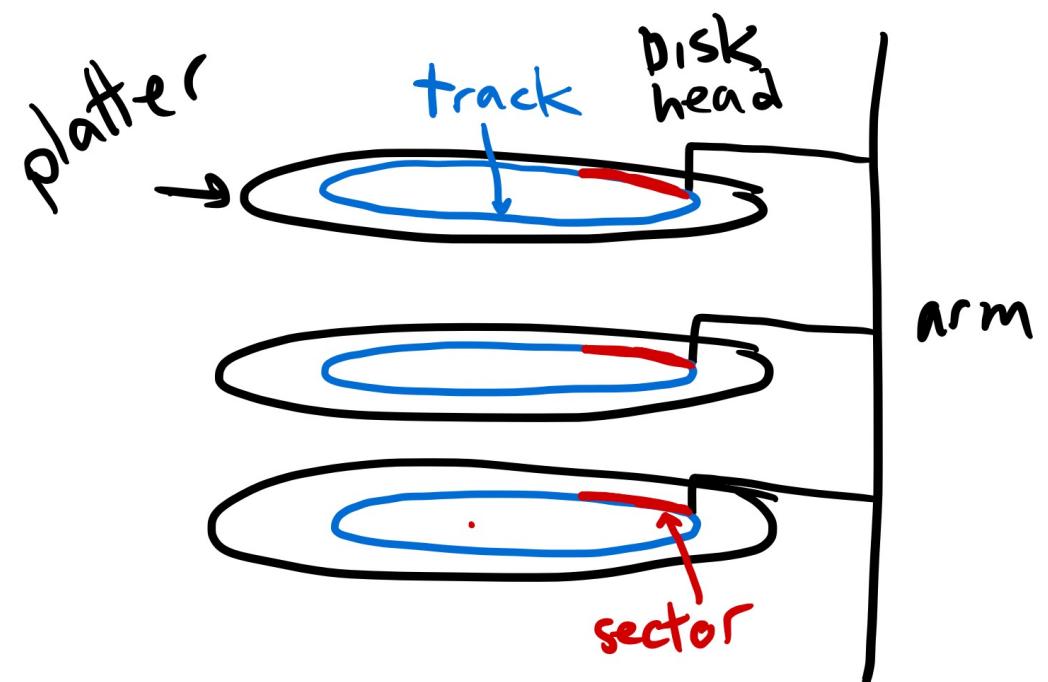
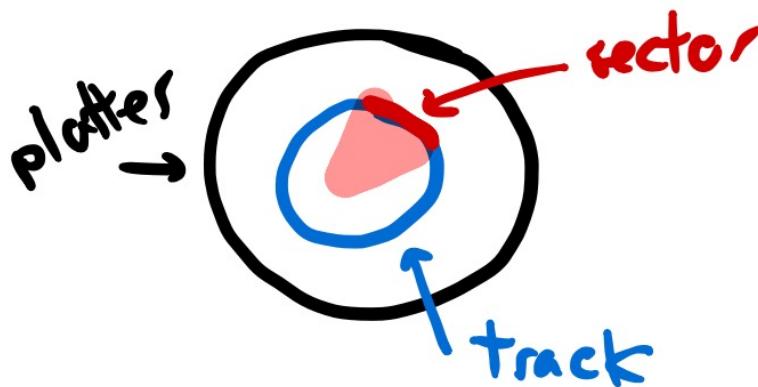


Spin speed: ~7200-15K RPM. 10K RPM ~ 6ms/rotation

Sector: multiples of pages.

More sectors further from center

Terminology: block = page in this class



Time to access (read or write) a disk block

seek time	2-4 msec avg
rotational delay	2-4 msec
transfer time	0.3 msec/64kb page

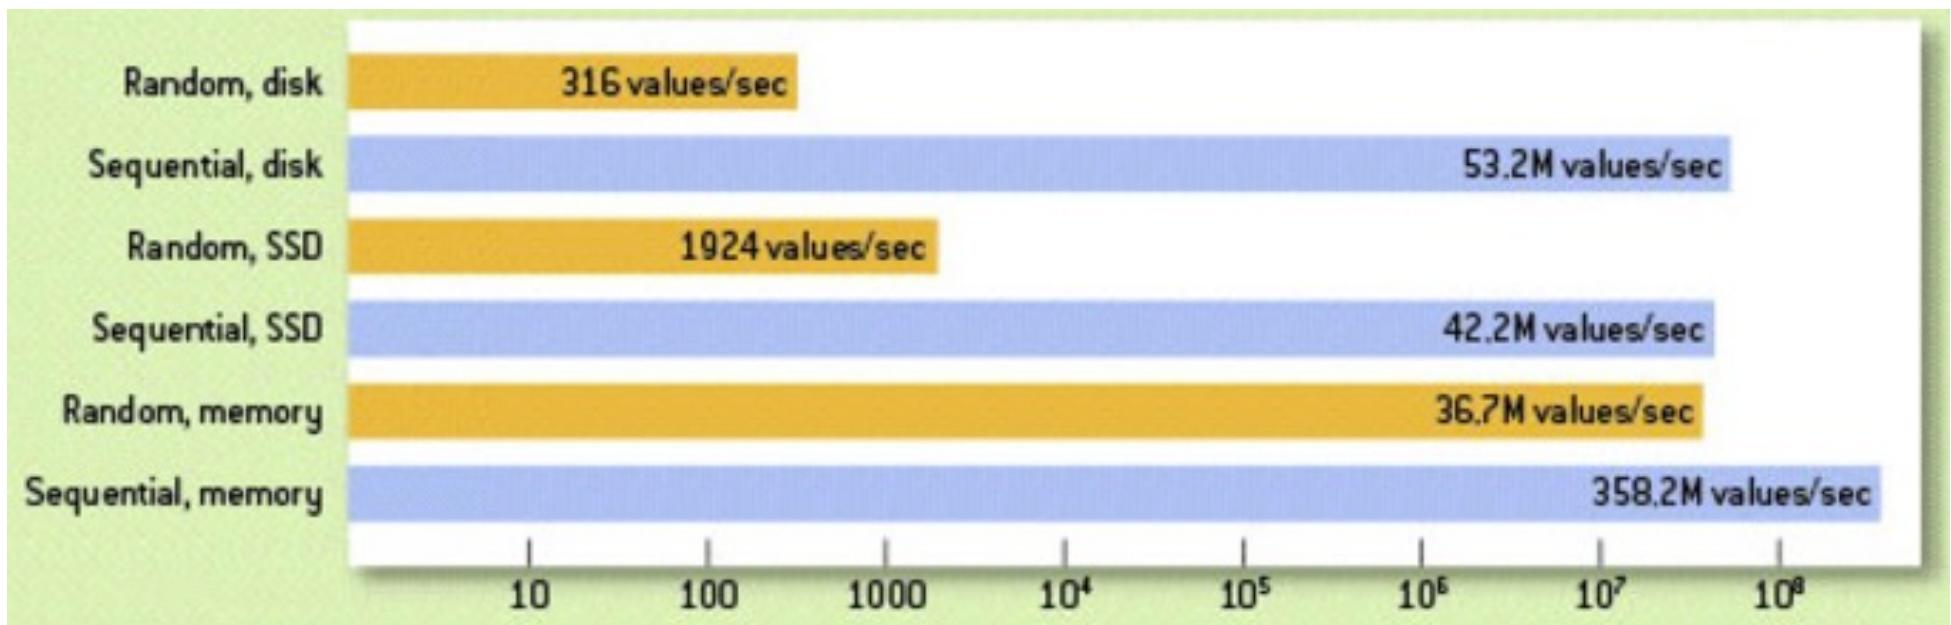
Throughput

read	~150 MB/sec
write	~50 MB/sec

Key: reduce seek and rotational delays

HW & SW approaches

# # of 4 byte values read per second



5 orders of magnitude

# SSDs

NAND memory

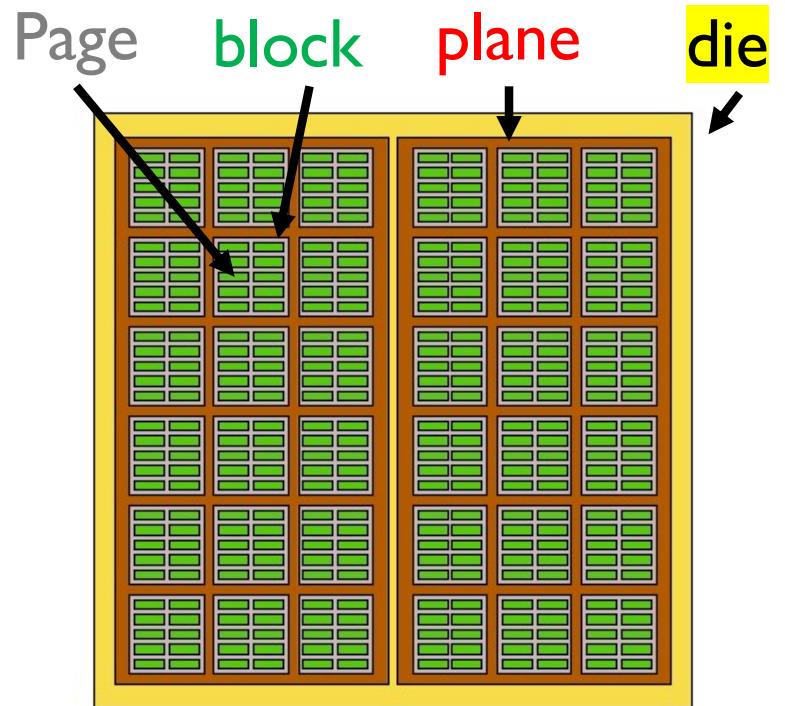
Small reads: 4-8k sized page

Big writes: 1-2MB block

2-3k writes before failure

wear leveling: distribute writes around

<https://www.anandtech.com/show/2738/5>



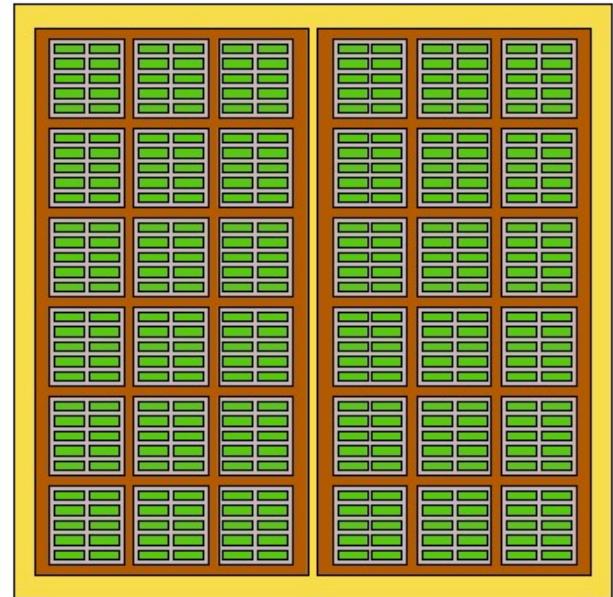
# SSDs

Reads fast

- single read time: 0.03ms
- random reads: 500MB/s
- sequential reads: 525MB/s

Writes less predictable

- single write time: 0.03ms
- random writes: 120MB/s
- sequential writes: 480MB/s



# What's Best? Depends on Application

Small databases: SSD/RAM

All global daily weather since 1929: 20GB

2000 US Census: 200GB

2009 english wikipedia: 14GB

Easily fits on an SSD or in RAM

Big databases: Disk

Sensors easily generate TBs of data/day

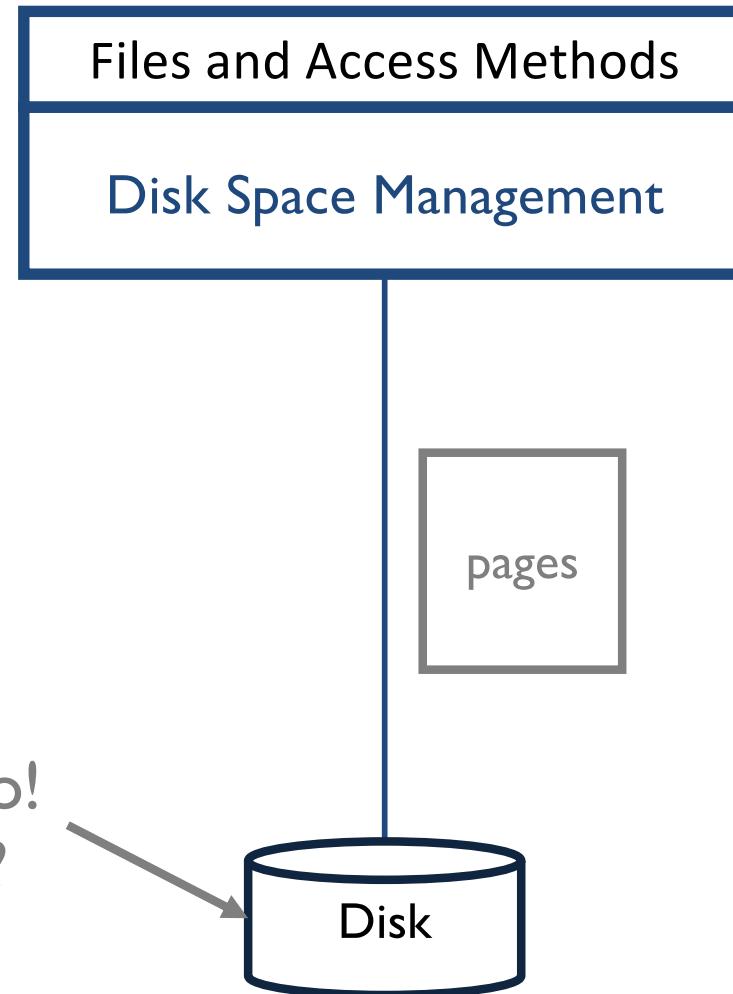
Boeing 787 generates  $\frac{1}{2}$  TB per flight

Disk has best cost-capacity ratio

SSDs help reduce read variance

AWS h1.16xlarge: 64 vCPU, 256GB ram, 8x2T hard drives

# Work from the bottom up



Like flying to Pluto!  
How to be faster?

# Strategies for Fast Data Access

(think about going to the store)

Big difference between random & sequential access

- Optimize for sequential accesses

**Amortize** sequentially read & write big chunks of bytes

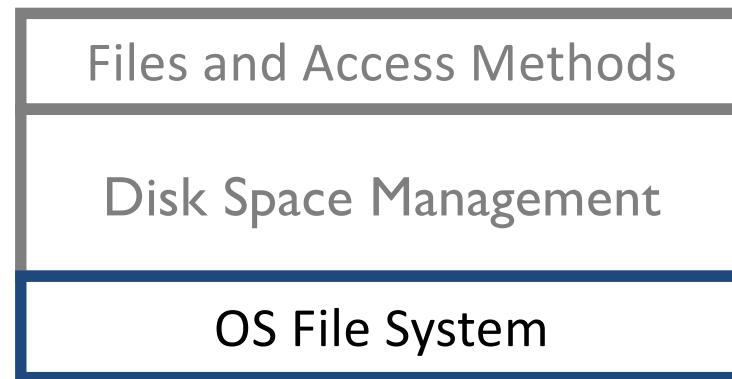
**Cache** popular blocks

**Pre-fetch** what you will need later

API

- read/write page
- read/write sequential pages
- notion of “next” page (upper layers can assume next is faster)

Big File



DBs don't directly talk to storage devices.

Go through OS file system

Big File

Files and Access Methods

Disk Space Management

Big File  
part 1

Big File  
part 2

Big File  
part 3

OS File System



OS File System



OS File System



# PostgreSQL database stored in files

```
ewu@dyn-160-39-248-197 ~/d/p/base> cd 33883/
ewu@dyn-160-39-248-197 ~/d/p/b/33883> ls
112          12526      2601      2612      2659      2696      2840_vm      3456_fsm      3602
113          12528      2601_fsm   2612_fsm   2660      2699      2841      3456_vm      3602_fsm
1247         12529      2601_vm    2612_vm    2661      2701      2995      3466      3602_vm
1247_fsm    12529_fsm  2602      2613      2662      2702      2995_vm    3466_vm    3603
1247_vm     12529_vm   2602_fsm   2613_vm    2663      2703      2996      3467      3603_fsm
1249         12531      2602_vm    2615      2664      2704      3079      3468      3603_vm
1249_fsm    12533      2603      2615_fsm  2665      2753      3079_fsm  3501      3604
1249_vm     12534      2603_fsm   2615_vm    2666      2753_fsm  3079_vm    3501_vm   3605
12504        12536      2603_vm    2616      2667      2753_vm   3080      3502      3606
12504_fsm   12538      2604      2616_fsm  2668      2754      3081      3503      3607
12504_vm    1255       2604_vm    2616_vm   2669      2755      3085      3534      3608
12506        1255_fsm  2605      2617      2670      2756      3118      3541      3609
12508        1255_vm   2605_fsm   2617_fsm  2673      2757      3118_vm   3541_fsm  3712
12509        1259       2605_vm    2617_vm   2674      2830      3119      3541_vm   3764
12509_fsm   1259_fsm  2606      2618      2675      2830_vm  3164      3542      3764_fsm
12509_vm    1259_vm   2606_fsm   2618_fsm  2678      2831      3256      3574      3764_vm
12511        1417       2606_vm    2618_vm   2679      2832      3256_vm  3575      3766
12513        1417_vm   2607      2619      2680      2832_vm  3257      3576      3767
12514        1418       2607_fsm   2619_fsm  2681      2833      3258      3576_vm  548
12514_fsm   1418_vm   2607_vm   2619_vm   2682      2834      3394      3596      549
12514_vm    174       2608      2620      2683      2834_vm  3394_fsm  3596_vm  826
12516        175       2608_fsm   2620_vm   2684      2835      3394_vm  3597      826_vm
12518        2187      2608_vm   2650      2685      2836      3395      3598      827
12519        2328      2609      2651      2686      2836_vm  34002     3598_vm  828
12519_fsm   2328_vm   2609_fsm   2652      2687      2837      34004     3599      PG_VERSION
12519_vm    2336      2609_vm   2653      2688      2838      34004_fsm 3600      pg_filenode.map
12521        2336_vm   2610      2654      2689      2838_fsm  34004_vm  3600_fsm  pg_internal.init
12523        2337      2610_fsm   2655      2690      2838_vm  34008     3600_vm
12524        2600      2610_vm   2656      2691      2839      34031     3601
12524_fsm   2600_fsm  2611      2657      2692      2840      3455      3601_fsm
12524_vm    2600_vm   2611_vm   2658      2693      2840_fsm  3456      3601_vm
ewu@dyn-160-39-248-197 ~/d/p/b/33883> █
```

# Administrivia

HW4 out

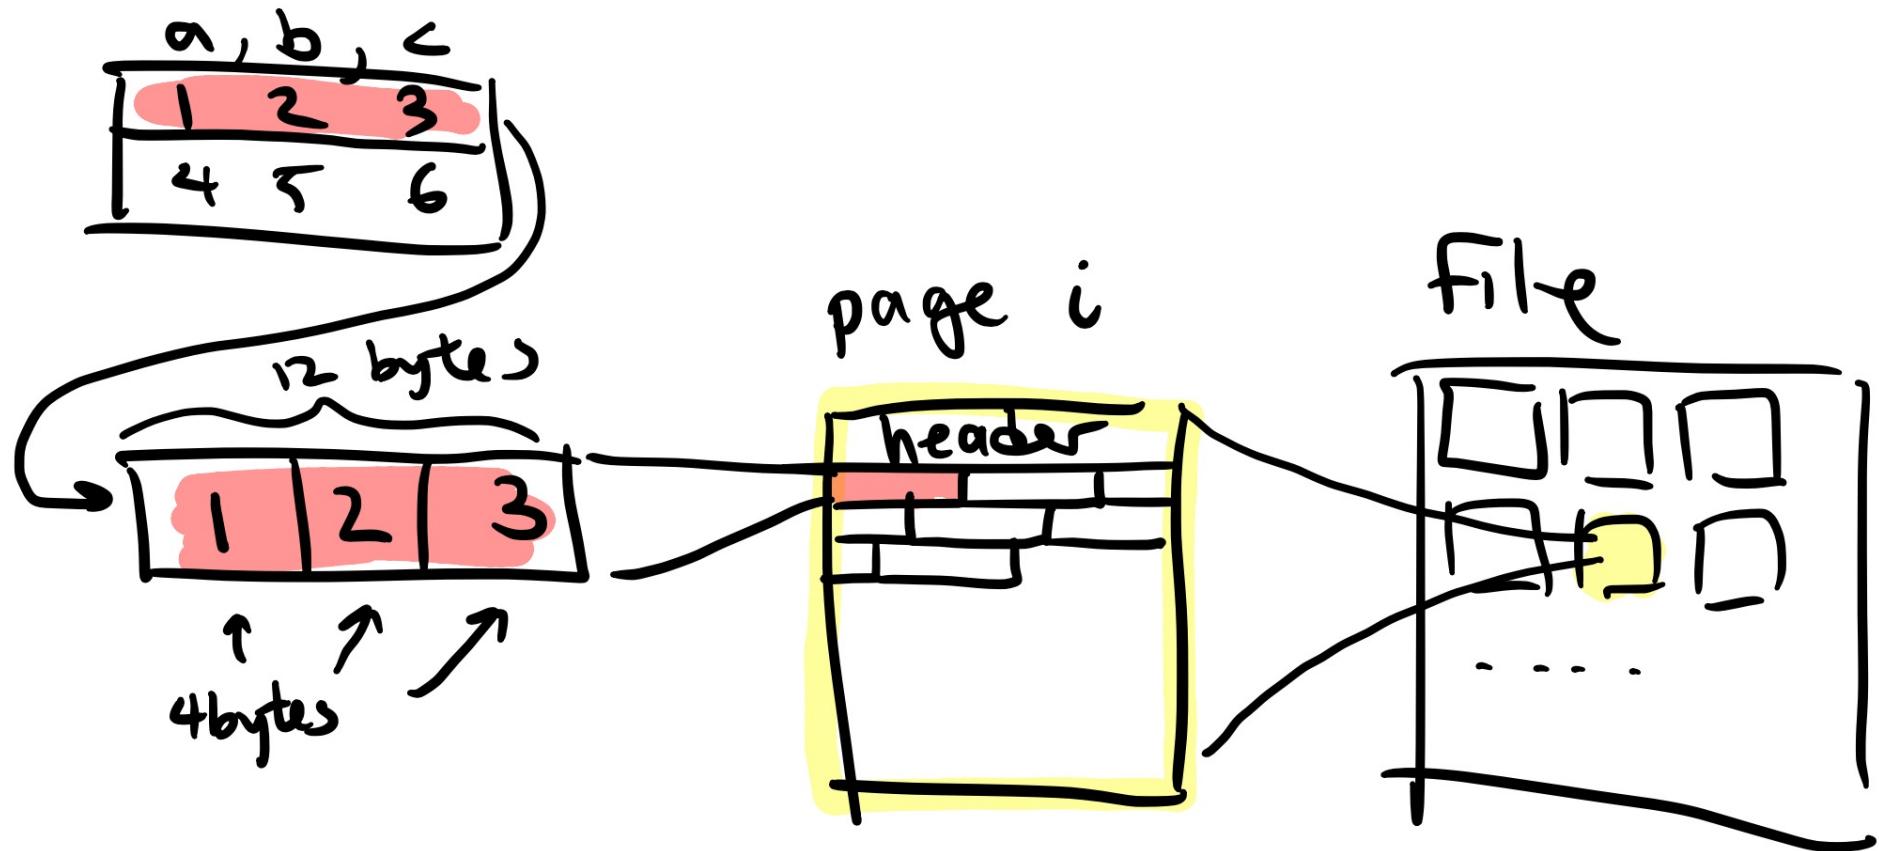
Project 2 out: **Start early!!!**

Both Due 11/29

**Maximum 3 late days!**

# Files

Logical relation



# Files

Higher layers want to talk in terms of records, and files of records

File: collection of pages

Minimum API:

- insert/delete/modify record

- lookup record\_id

- scan all records

Page: collection of records

- typically fixed page sizes (8 or 64kb in PostgreSQL)

These are logical.

# File Organizations Overview

What do pages store, and how to organize pages?

- Unordered heap file
- Heap file
- B+ tree index
- Hash index

# Units that we'll care about

Ignore CPU cost

Ignore RAM cost

B # *data* pages on disk for relation

R # records per data page

D avg time to read/write data page to/from disk

Simplifies life when computing costs

OK to not be exactly correct

# Unordered Heap Files

Collection of records (no order)

As we add records, pages allocated

As we remove records, pages removed

To support record level ops, need to track:

pages in file

free space on pages

records on page

Ok, let's design that

# Super Naïve Design

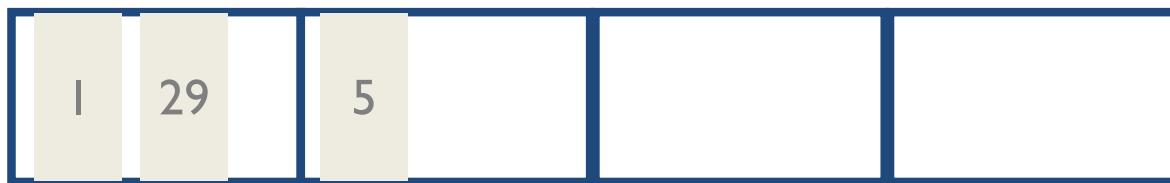


Big array  
of bytes

Heap is a big array of bytes.

First split into array of pages

# Super Naïve Design



Array of  
pages

Insert



Each each page from disk sequentially

Check if there's an open spot

Insert 4 into open spot.

# Super Naïve Design

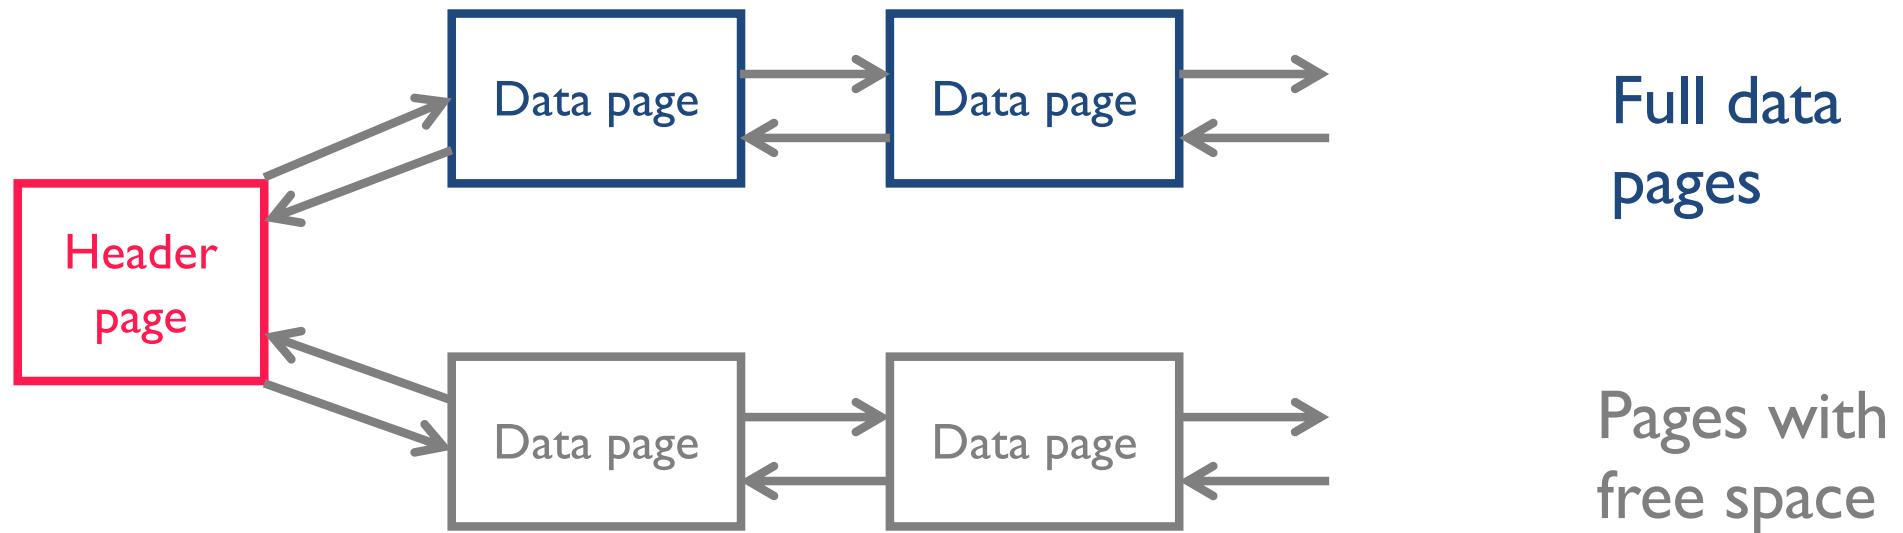


Array of  
pages

Some problems include

- Slow to know what pages are empty, full, have space
- Slow to find a particular value
- Fragmentation

# Heap File



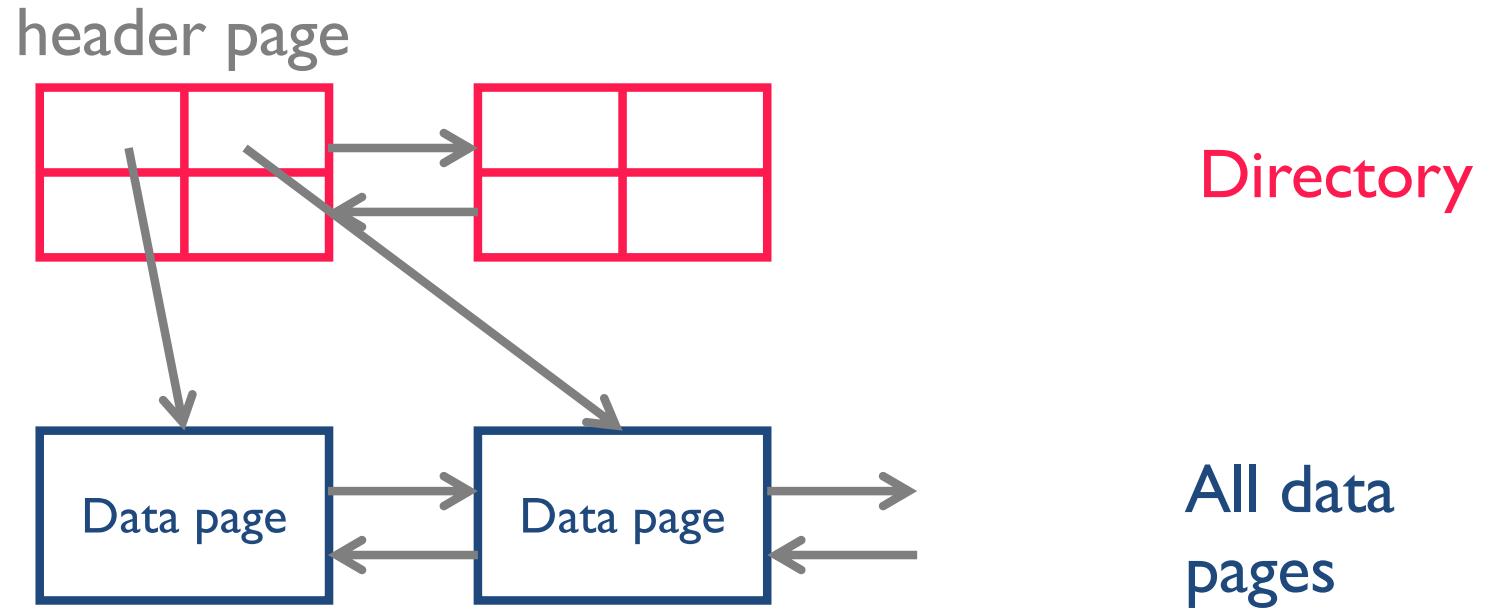
Header page info stored in catalog

Data page contains: 2 pointers, free space, data

Need to scan pages to answer any query

Which page has enough free space for 100 bytes?

# Use a directory



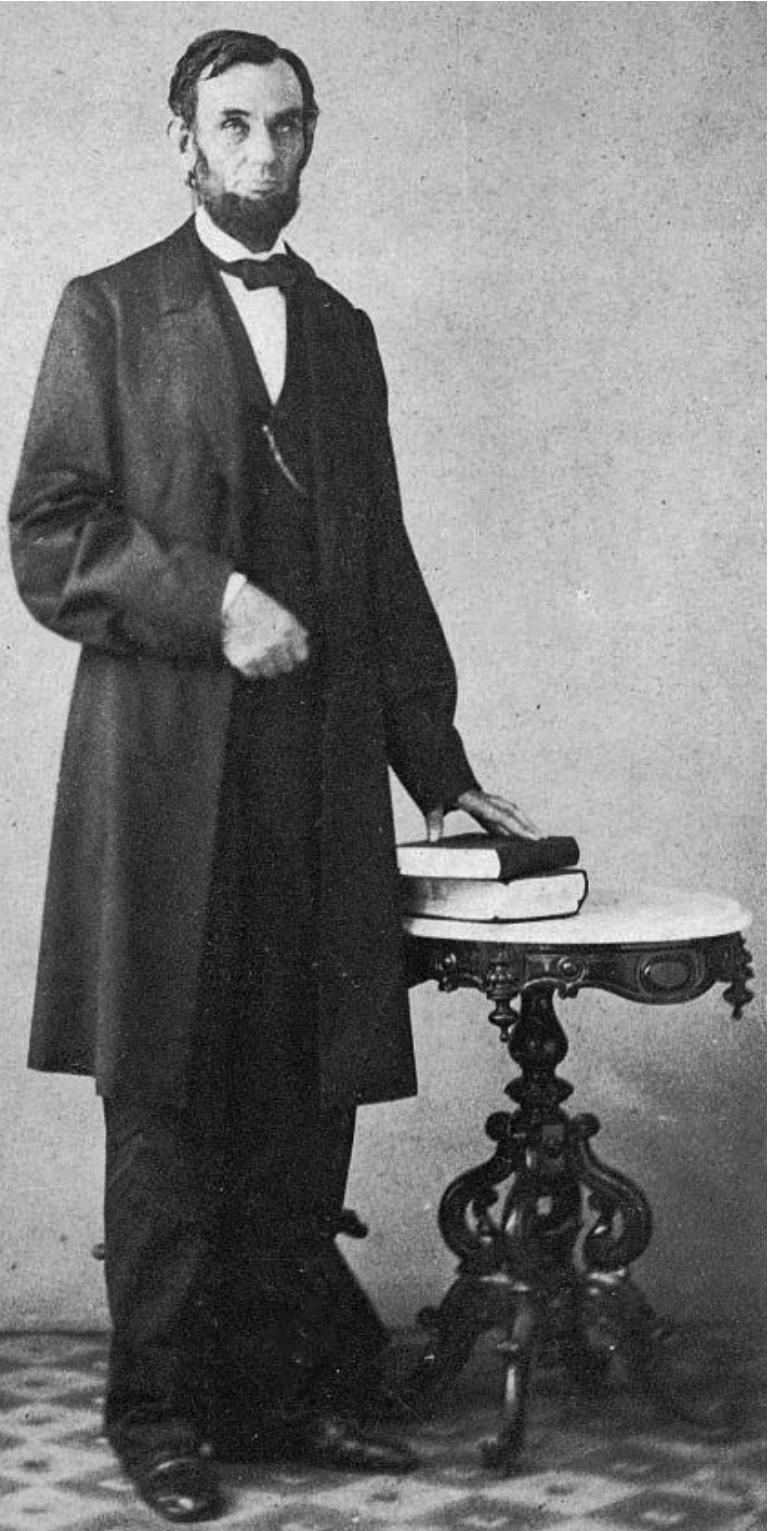
Directory entries track # free bytes on data pages

Directory is collection of pages

# Indexes

“If I had 8 hours to  
chop down a tree,  
I'd spend 6 sharpening  
my ax.”

*Abraham Lincoln*



# Indexes

Heap files answer any query via a sequential scan, but...

Queries use *qualifications* (predicates)

find students where class = “CS”

find students with age > 10

Indexes: file structures for value-based queries

B+-tree index (~1970s)

Hash index

Overview! Details in 4112

# Indexes

Defined with respect to a *search key*

don't confuse with candidate keys!!

Faster access for WHERE clauses w/ search key

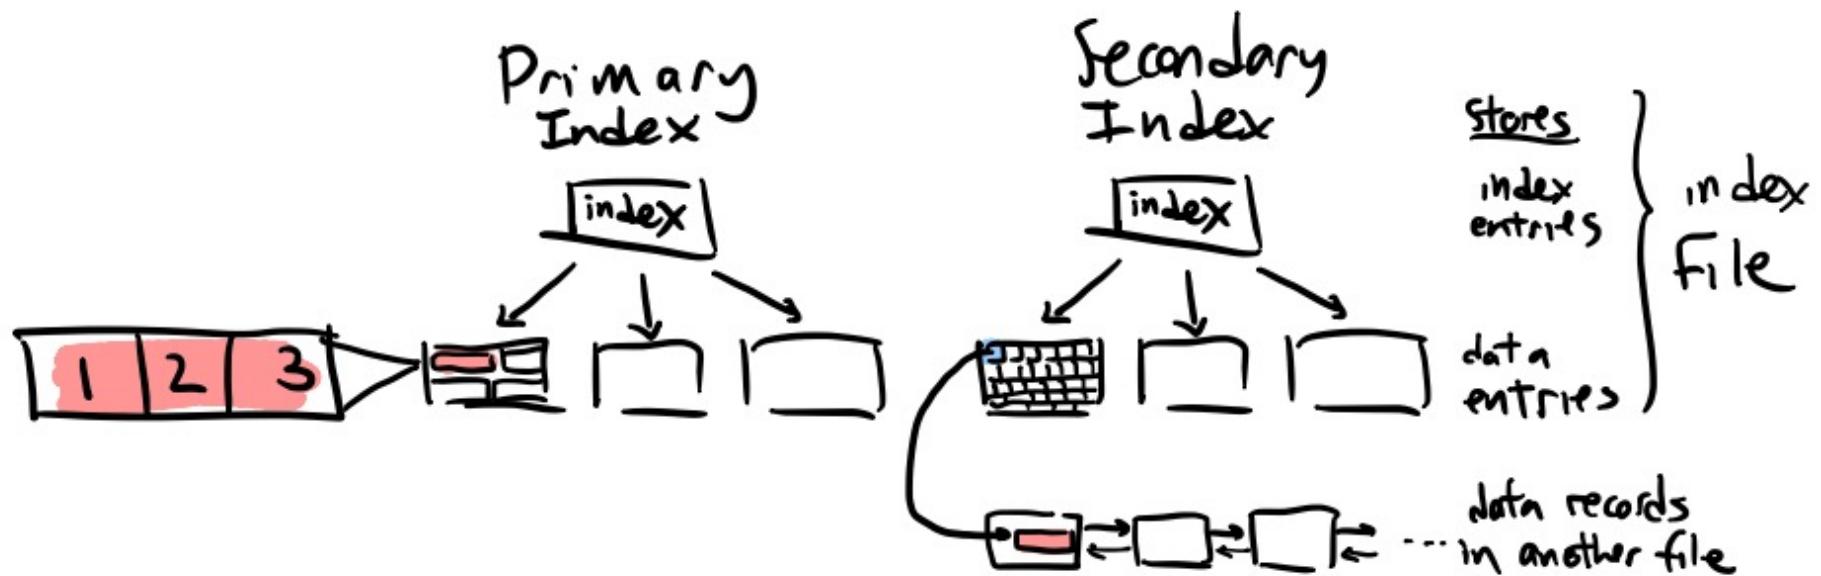
```
CREATE INDEX idx1 ON users USING btree (sid)
```

```
CREATE INDEX idx2 ON users USING hash (sid)
```

```
CREATE INDEX idx3 ON users USING btree (age,name)
```

You will play around with indexes in HW4

# Primary vs Secondary Index Files



## Primary

Data entries contain actual tuples

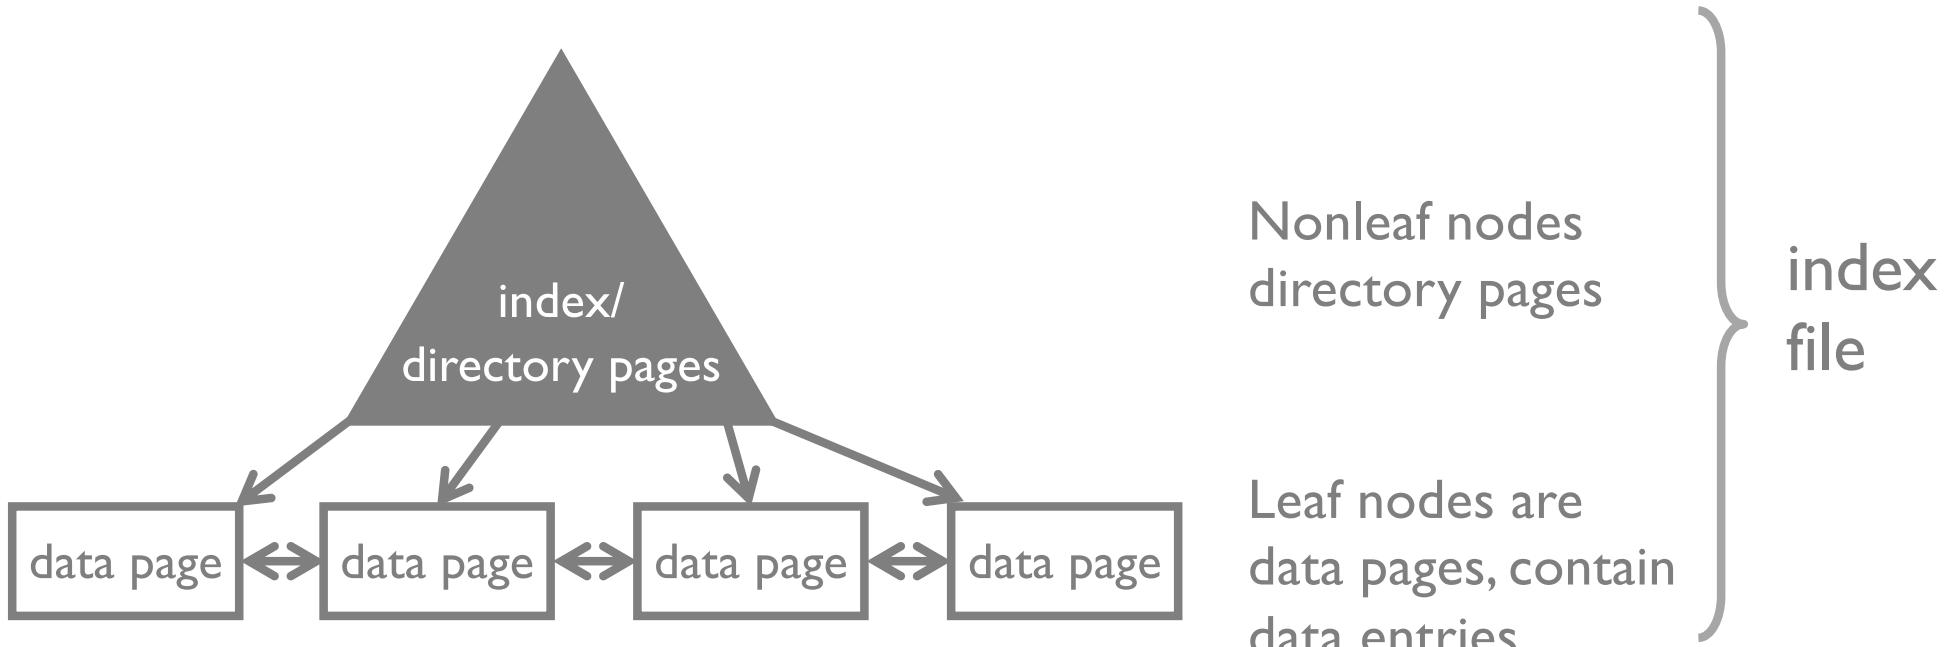
Pros: directly access data.

## Secondary

Data entries contain <search key val, rid> pointers into another file

Pros: index is more compact

# B+ Tree Index: disk-optimized index



**Node = Page**

Supports equality and range queries on search key

Self balancing (path to any leaf is almost the same)

Leaf nodes are connected via doubly linked list

Height is with respect to directory pages (the gray part of the triangle)

# B+ Tree on (age)

**Non-leaf Directory pages**

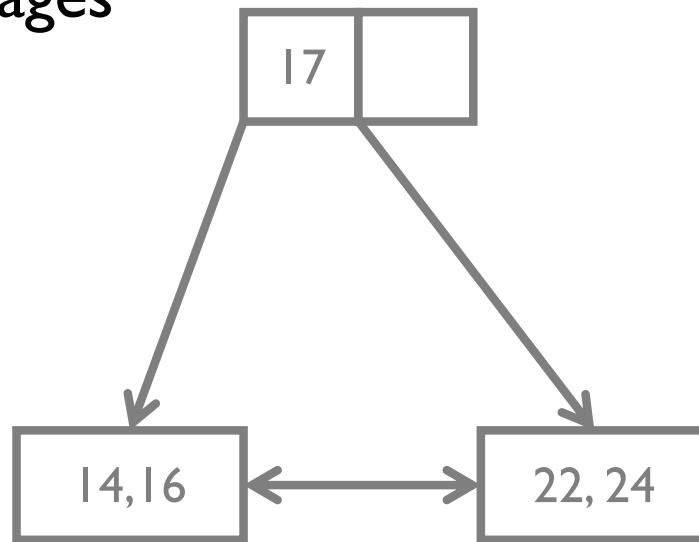
m index entries

m+1 pointers

**Leaf Data pages**

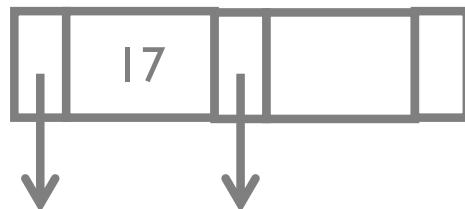
data entries/tuples

At each level in the tree, index & data page contents are sorted by search key



Query:    **SELECT \* WHERE age= 14**

directory page

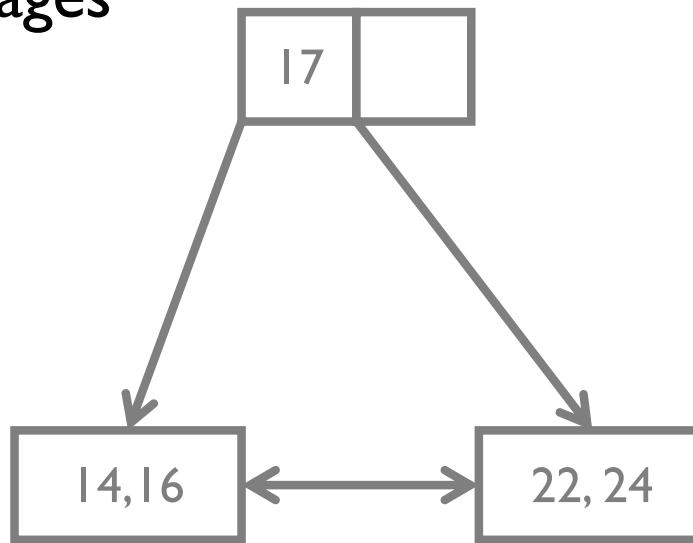


# Index Only Queries: B+ Tree on (age)

Non-leaf Directory pages

m index entries

m+l pointers



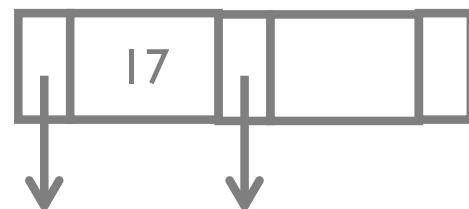
At each level in the tree, index & data page contents are sorted by search key

Leaf Data pages

data entries/tuples

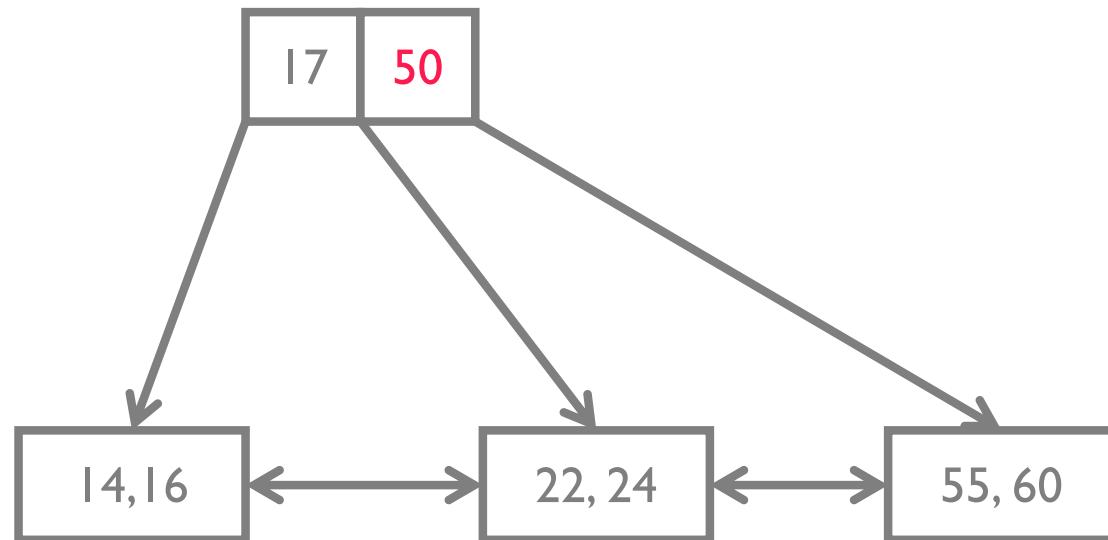
Query:    **SELECT age WHERE age > 14**  
**(index only!)**

directory page



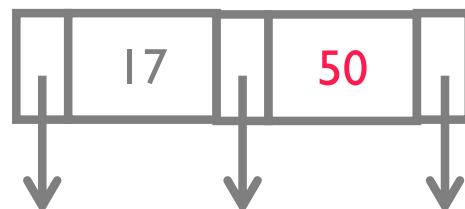
# B+ Tree on (age)

Note: 50 not a data entry



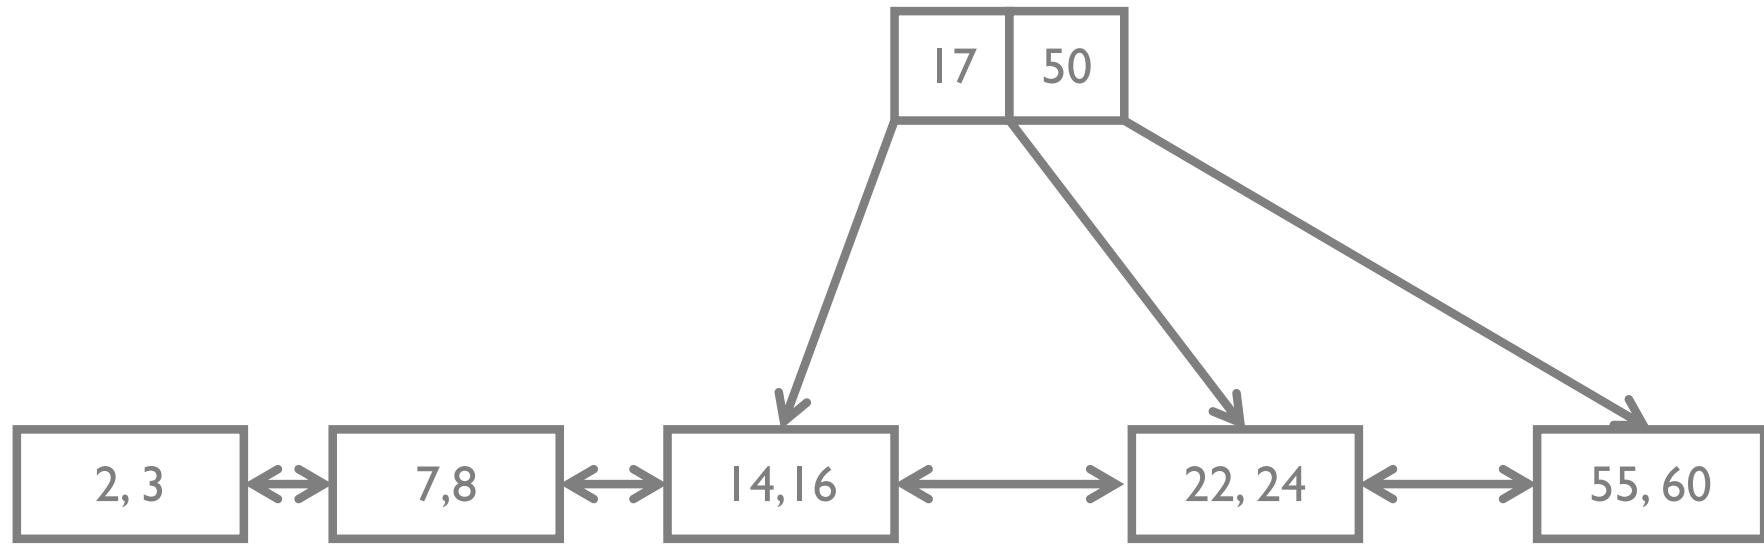
Query: **SELECT \* WHERE age = 55**

directory page



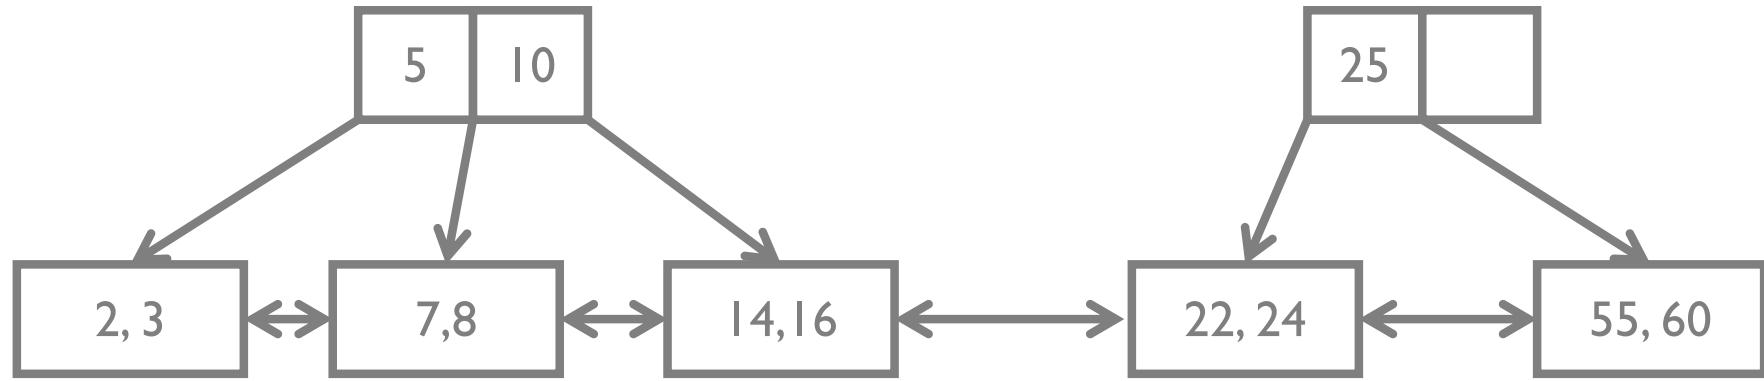
In this class for cost estimation, we will assume number of pointers same as number of directory entries

# B+ Tree on (age)



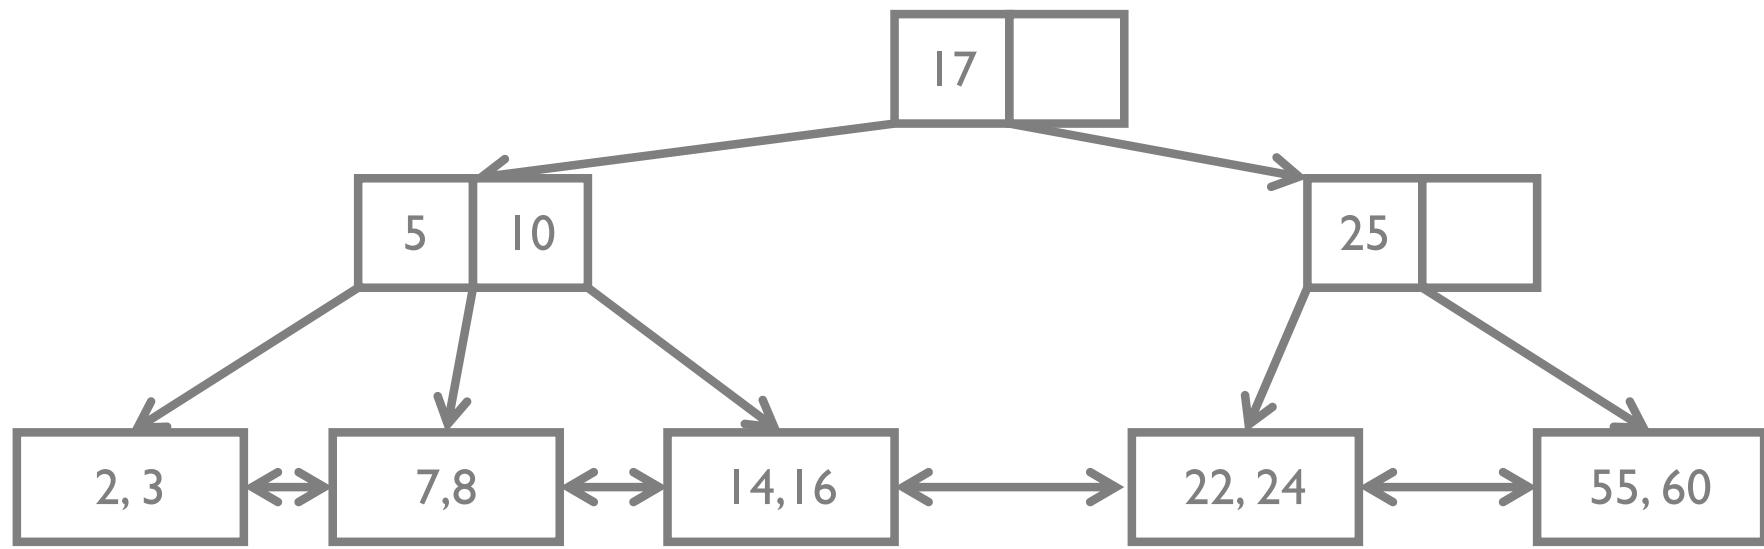
When directory page is full, can't add more pointers  
Split directory page and reorganize tree  
(don't need to know details of how splits are done)

# B+ Tree on (age)



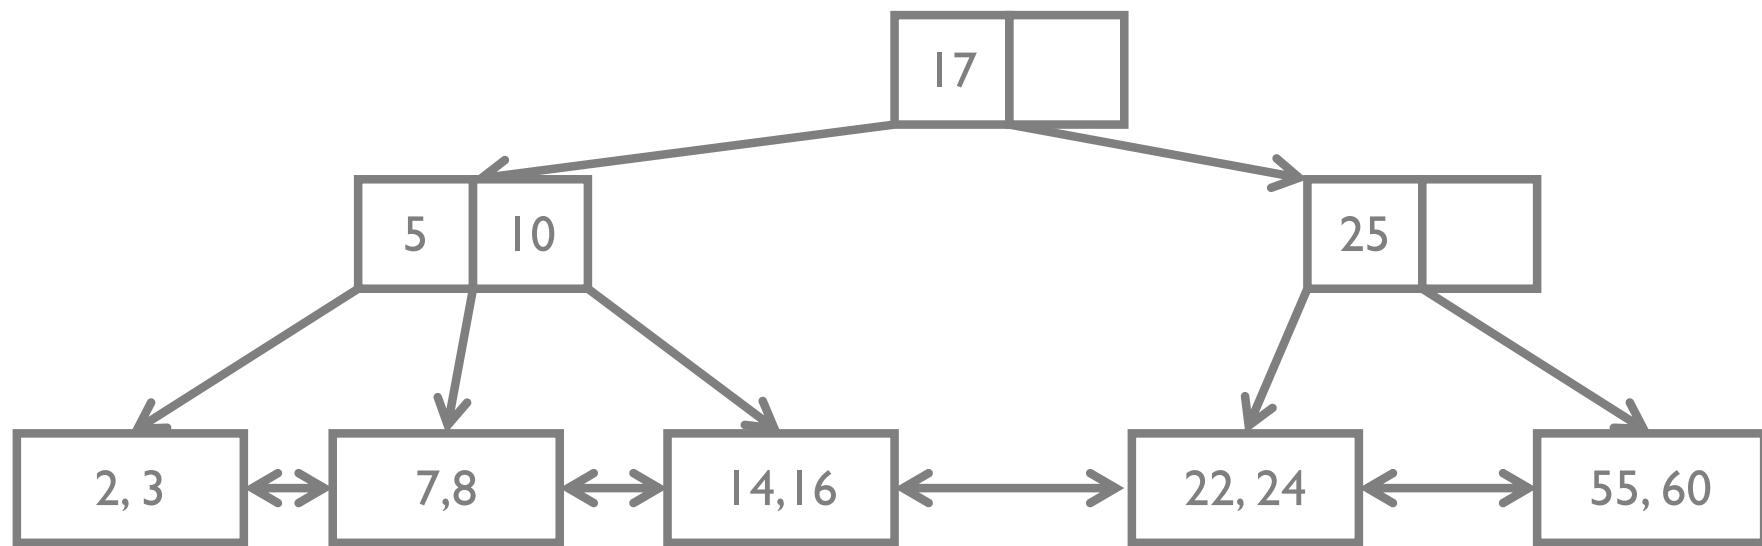
When directory page is full, can't add more pointers  
Split directory page and reorganize tree  
(don't need to know details of how splits are done)

# B+ Tree on (age)



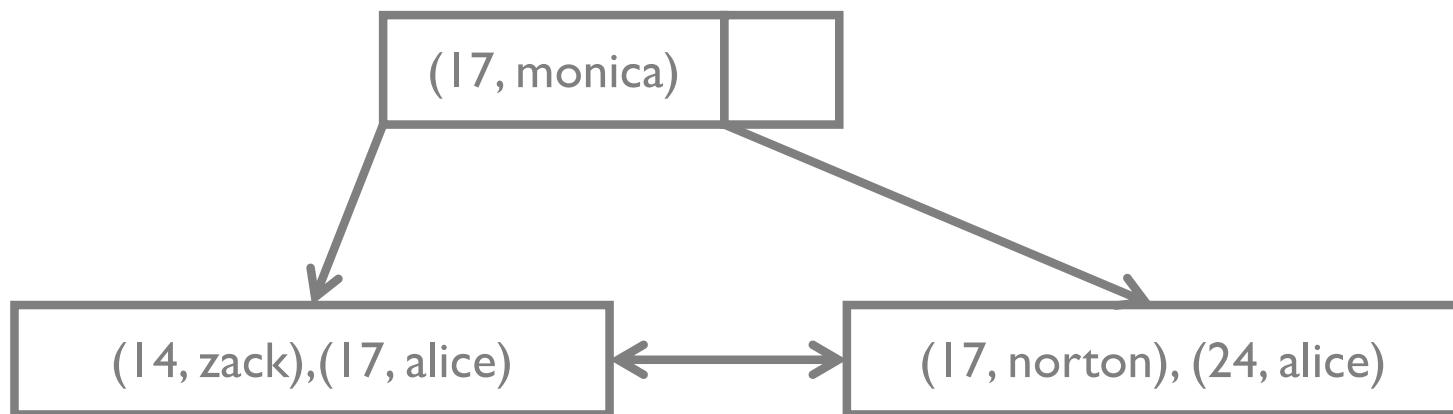
When directory page is full, can't add more pointers  
Split directory page and reorganize tree  
(don't need to know details of how splits are done)

# B+ Tree on (age)



Query: **SELECT \* WHERE age > 20**

# Composite Keys: B+ Tree on (age, name)



How do the following queries use the index on (age, name)?

SELECT age WHERE age = 14

SELECT \* WHERE age < 18 AND name < 'monica'

SELECT age WHERE name = 'bobby'

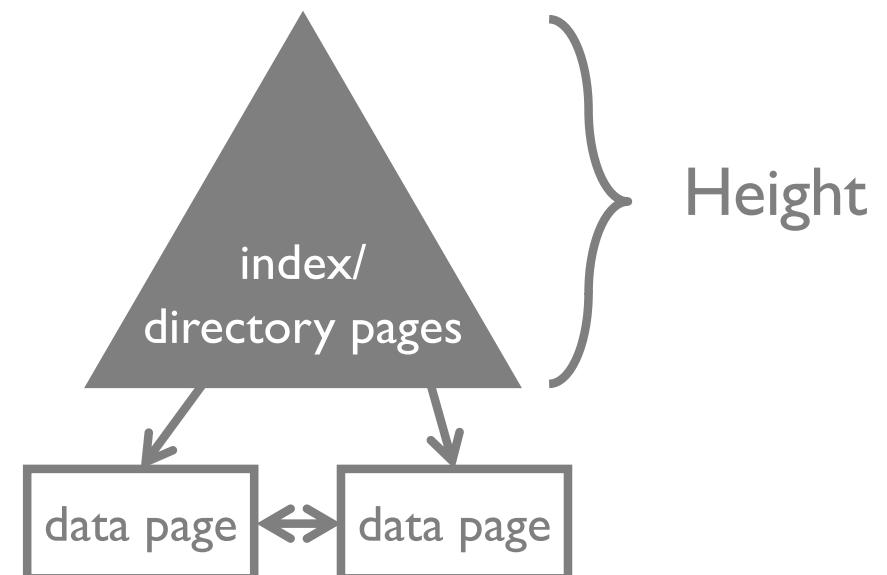
# Terminology



Keep free space for fast inserts  
Reduce “fill factor” parameter



Fanout (“branching factor”)  
We assume same as # directory entries



# Example using 8kb pages

How many entries in the index's data pages?

fill-factor: ~66%

~300 entries per directory page

height 2:  $300^3 \sim 27$  Million entries assuming 300 tuples/pg

height 3:  $300^4 \sim 8.1$  Billion entries assuming 300 tuples/pg

Top levels often in memory

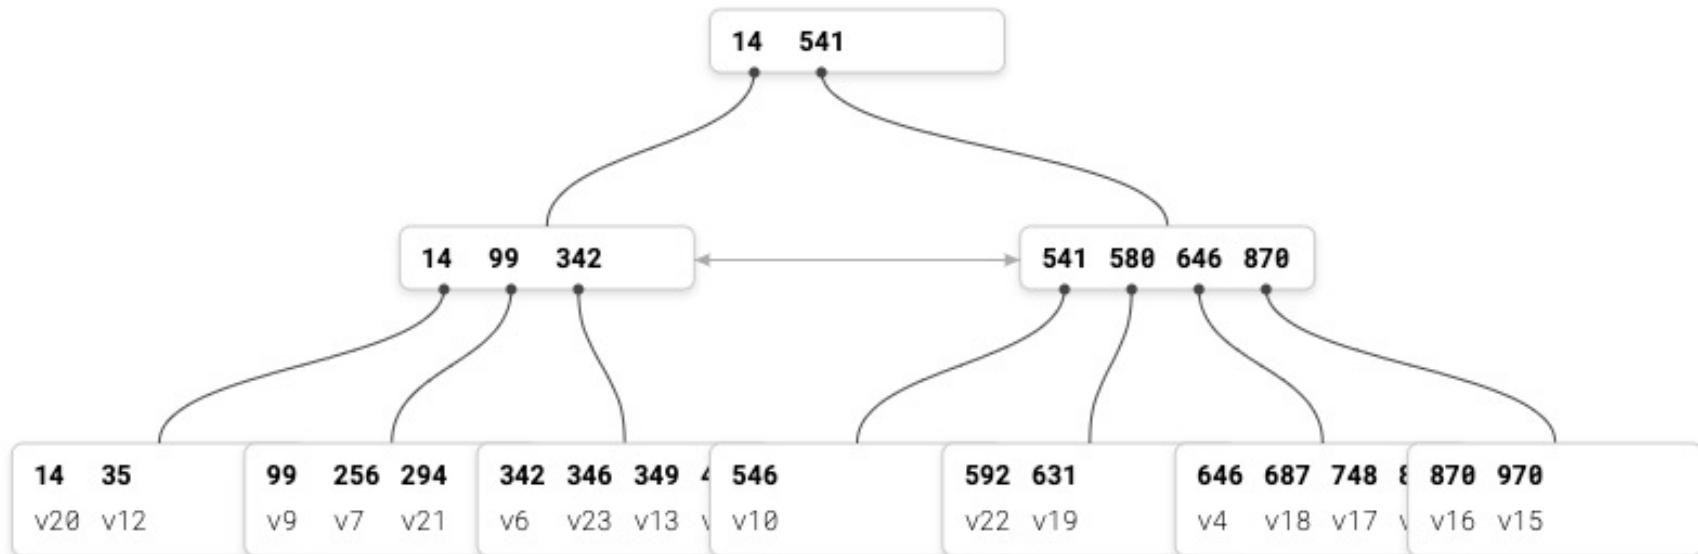
height 2 only 300 pages ~2.4MB

height 3 only 90k pages ~750MB

# For Real?

problem. We have run TB range databases on Raspberry PI with 4GB of RAM and hammered that in benchmarks (far exceeding the memory capacity). The interesting thing here is that the B+Tree nature means that the upper tiers of the tree were already in memory, so we mostly ended up with a single page fault per request.

<https://bplustree.app/>



# Hash Index on age

Hash function

$$h(v) = v \% 3$$

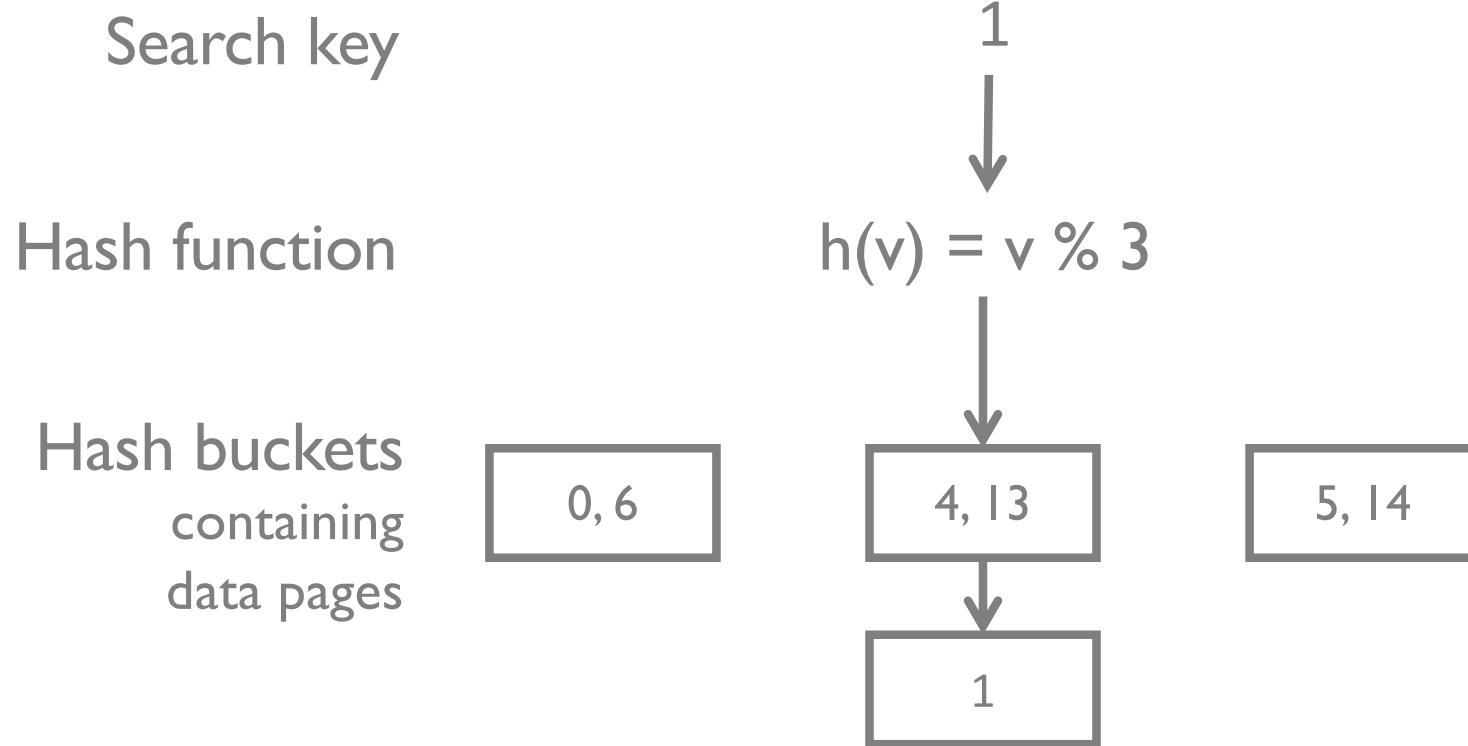
Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

# INSERT Hash Index on age



# INSERT Hash Index on age

Search key

11

Hash function

$$h(v) = v \% 3$$

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

1



# INSERT Hash Index on age

Search key

11

$$h(v) = v \% 3$$

Hash function

Hash buckets  
containing  
data pages

0, 6

4, 13

5, 14

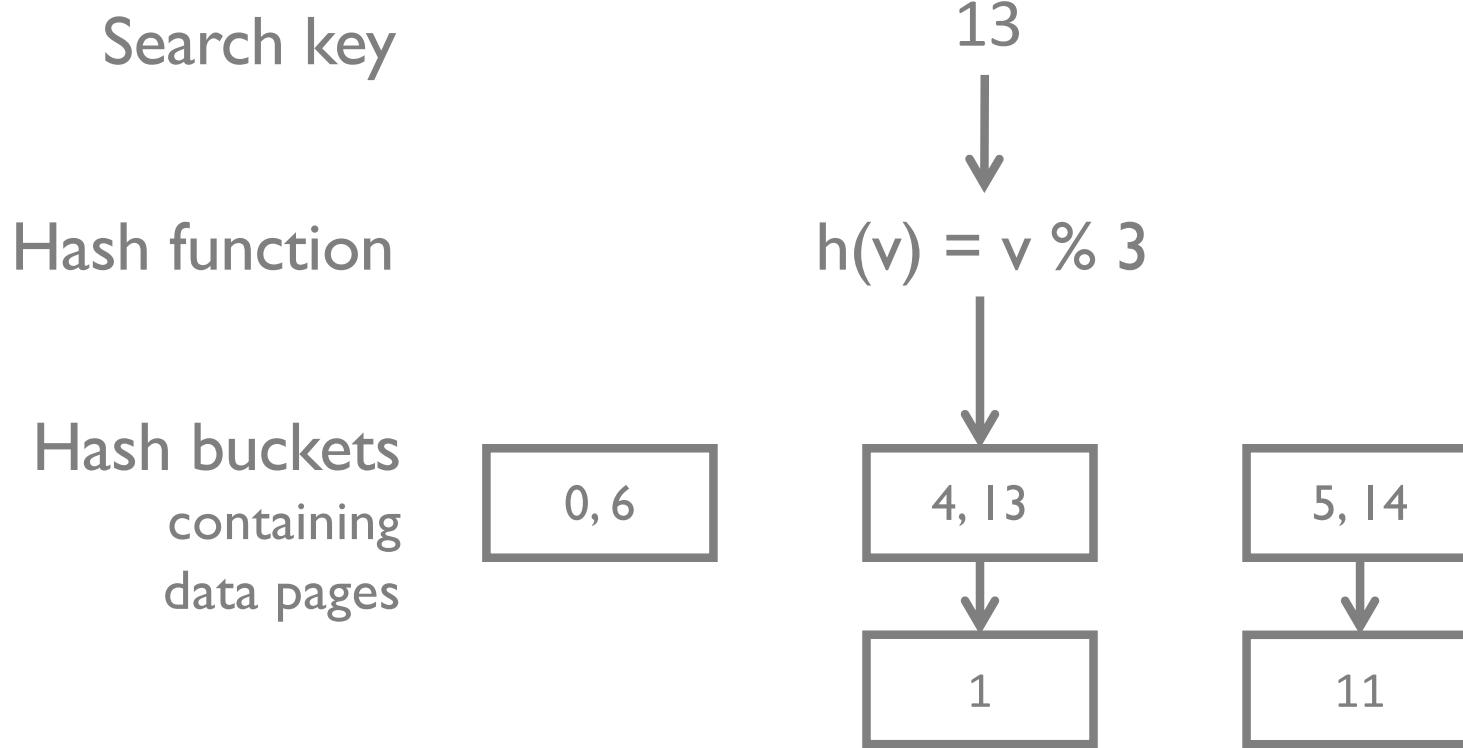
1

11

11



# SEARCH Hash Index on age



Good for equality selections

Index = data pages + overflow data pages

Hash function  $h(v)$  takes as input the search key

# Costs

Three file types

Heap, B+ Tree, Hash

Indexes can be primary or secondary

Operations we care about

Scan all data    `SELECT * FROM R`

Equality            `SELECT * FROM R WHERE x = 1`

Range              `SELECT * FROM R WHERE x > 10 and x < 50`

Insert tuple

Delete tuple    `DELETE WHERE ...`

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything				
Equality				
Range				
Insert				
Delete				

The above table's calculations will be based on assumptions about the fill factors, directory and data entry sizes, page sizes, etc. They estimate the cost in expectation.

- B** # data pages
- D** time to read/write page
- M** # pages in range query

The assumptions may vary depending on the system and configuration.

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD			
Equality	0.5BD			
Range	BD			
Insert	2D			
Delete	Search + D			

## Heap File

equality on a key. How many results?

- B # data pages
- D time to read/write page
- M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD		
Equality	0.5BD	$D(\log_2 B)$		
Range	BD	$D(\log_2 B + M)$		
Insert	2D	Search + BD		
Delete	Search + D	Search + BD		

## Heap File

equality on a key. How many results?

## Sorted File

files compacted after deletion

- B # data pages
- D time to read/write page
- M # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + I)$	
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	
Insert	2D	Search + BD	$D(\log_{80} B + 2)$	
Delete	Search + D	Search + BD	$D(\log_{80} B + 2)$	

## Heap File

equality on a key. How many results?

## Sorted File

assume: files compacted after deletion

## B+ Tree

100 entries/directory page

80% fill factor

**B** # data pages

**D** time to read/write page

**M** # pages in range query

	Heap File	Sorted Heap	B+ Tree	Hash
Scan everything	BD	BD	1.25BD	1.25BD
Equality	0.5BD	$D(\log_2 B)$	$D(\log_{80} B + I)$	D
Range	BD	$D(\log_2 B + M)$	$D(\log_{80} B + M)$	1.25BD
Insert	2D	Search + BD	$D(\log_{80} B + 2)$	2D
Delete	Search + D	Search + BD	$D(\log_{80} B + 2)$	2D

## Heap File

equality on a key. How many results?

## Sorted File

assume: files compacted after deletion

**B** # data pages  
**D** time to read/write page  
**M** # pages in range query

## B+ Tree

100 entries/directory page

80% fill factor

## Hash index

80% fill factor when computing scans

assumes no overflow when computing equality/insert/delete.

# Administrivia

## Extra credit 2 is out (2%)

- analyze Kaggle competition data using SQL
- compare analysis using normalized and denormalized schemas
- help do research (optional)

### ▶ Consent Form

► **Purpose of this study:** We hope to use assignment submissions as part of a research study. The purpose of this study is to understand how easy or difficult it is to write correct analysis queries over normalized and denormalized databases. This will help us understand challenges in database decomposition, and inform the development of simpler and more correct querying interfaces. If you have any questions about the study, please reach out to Zachary Huang ([zh2408@columbia.edu](mailto:zh2408@columbia.edu)).

**Freedom to withdraw:** Your choice to participate in this study is completely voluntary and does not affect your extra credit grade. If you want to withdraw your consent in the future, please contact Zachary Huang. We will destroy any data we hold from you.

**Privacy & Data Security:** We will ask you to write SQL queries using normalized and denormalized databases, and fill out surveys. We do not collect any other personal information and the data are anonymous. We will carefully examine the words and exclude any details that may contain identifiable information. Identifying information involved during recruitment (uni) will not be stored for the study. We may publish research reports that include your comments and actions, but they will be anonymous.

If you are below 18, we will not use your data.

**If you are over 18 and are willing to participate in the study, please answer yes. If you are below 18 or decline to participate in the study, please answer no.**

**consent:** yes

# ucdsfair.com



Vagelos Computational  
Science Center  
BARNARD COLLEGE

CS<sup>†</sup>  
@CU COMPUTER SCIENCE

[Home](#) [Sponsor](#)  
[Exhibitors](#)

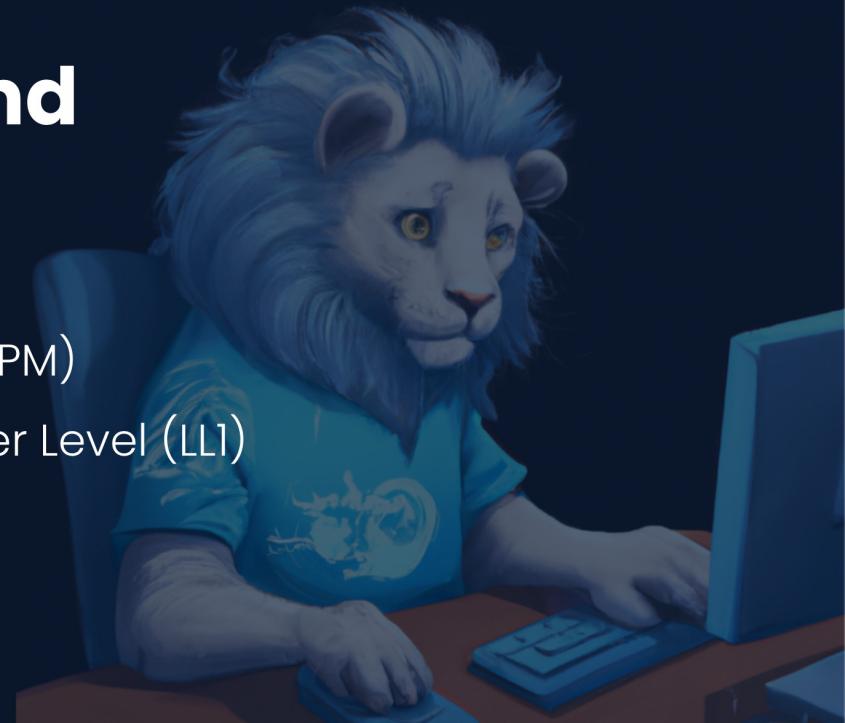
[Register Here](#)

## Undergraduate Computer and Data Science Research Fair

**Event Date:** Friday, November 11, 2022 (2:00 PM - 6:00 PM)

**Location:** Barnard Diana Center, The Event Oval, Lower Level (LL1)

[Register Here](#)



# Where do B, D, M come from?

Estimated from more basic info

Assuming (we will be clear about this)

- fanout = number of directory entries
- pointer in secondary index same size as directory entry

Given:

- p: page size
- r: record size
- d: directory entry size
- f: fill factor
- n: # records

Estimate for primary and secondary index:

- size
- height
- access cost

# Primary B+ Index

Page Size	$p = 100$	records/page	$p / r = 10$
Record Size	$r = 10$	direntries/page	$p / d = 20$
Dir entry Size	$d = 5$	fanout:	20
Fill Factor	$f = 100\%$	# data pages	$n/(p/r) = 800$
# Records	$n = 8000$	height	$\log_{20}800 = 3$

Cost to look up a single record is  
3 for directory pages + 1 data page

# Secondary B+ Index

Page Size	$p = 100$	records/page	$p / r = 10$
Record Size	$r = 10$	direntries/page	$p / d = 20$
Dir entry Size	$d = 5$	fanout:	20
Fill Factor	$f = 100\%$	# data pages	$n/(p/d) = 400$
# Records	$n = 8000$	height	2

Cost to look up a single record is  
2 for directory pages + 1 data page + 1 pointer lookup

# How to pick?

Depends on your queries (workload)

Which relations?

Which attributes?

Which types of predicates (=, <,>)

*Selectivity*

Insert/delete/update queries? how many?

# How to choose indexes?

## Considerations

which relations should have indexes?

on what attributes?

how many indexes?

what type of index (hash/tree)?

# Naïve Algorithm

get query workload

group queries by type

for each query type in order of importance

    calculate best cost using current indexes

    if new index IDX will further reduce cost

        create IDX

Why not create every index?

update queries slowed down (upkeep costs)

takes up space

# High level guidelines

Check the WHERE clauses

- attributes in WHERE are search/index keys

- equality predicate → hash or tree index

- range predicate → tree index

Multi-attribute search keys supported

- order of attributes matters for range queries

- may enable queries that don't look at data pages (*index-only*)

# Summary

Design depends on economics, access cost ratios

Disk still dominant wrt cost/capacity ratio

Many physical layouts for files

same APIs, difference performance

remember physical independence

## Indexes

Structures to speed up read queries

Multiple indexes possible

Decision depends on workload

# Things to Know

How a hard drive works and its major performance characteristics

The storage hierarchy & differences between RAM, SSD, Hard drives

What files, pages, and records are, and how different than UNIX model

Heap File data structure

B+ tree and Hash indexes

Performance characteristics of different file organizations

Given statistics, figure out directory size, index height, access cost